



Experimenting with Reaction Systems using Graph Transformation and GROOVE

Roberto Bruni¹ · Arend Rensink²

Accepted: 27 July 2025
© The Author(s) 2025

Abstract

We explore the capabilities of GROOVE, a state-of-the-art toolset based on graph transformation systems, to perform different kinds of analyses of Reaction Systems, ranging from reachability and causal analysis to model checking. Our results are encouraging, as in the presence of large state spaces GROOVE improves the time required for both reachability and causal analyses by an order of magnitude, compared to other available tools. From the point of view of GROOVE, the implementation of Reaction Systems provided some interesting insights on the most convenient way to model certain computational requirements through negative and nested application conditions.

Keywords Reaction Systems · Graph Transformation · GROOVE · Causal analysis · Formal methods

1 Introduction

We explore the capabilities of a state-of-the-art toolset based on graph transformation to perform reachability, causal analysis and verification of complex systems modelled using Reaction Systems.

Reaction systems (RSs) (Ehrenfeucht and Rozenberg 2007) are a computational model inspired by the functioning of biochemical reactions within living cells. RSs focus on the interaction of entities through a set of reactions. Each reaction relies on some reactants, inhibitors, and products to mimic two fundamental mechanisms found in nature: facilitation and inhibition. At each time instant, the next state of the system is determined by the products of all enabled reactions plus some additional entities that are possibly provided by the environment. Unlike traditional models of concurrency, like Petri nets, the theory of RSs is based on three principles: *no permanency*: any entity vanishes unless it is sustained by a reaction; *no competition*: an entity is either available for

all reactions, or it is not available at all; and *no counting*: the exact concentration level of available entities is ignored. Moreover, due to the use of inhibitors, RS can exhibit non-monotonic behaviour, in the sense that what can be done with fewer resources is not necessarily replicable with more resources.

While *closed* RSs evolve in isolation, *interactive processes* are dealt with by providing suitable environments that provide a sequence of stimuli at each step: these are sets of entities that can be used to trigger or inhibit some reactions. A common example involves using contexts to analyse how drug administration affects organisms that are modelled as sets of reactions.

Since their introduction, RSs have been successfully applied to the analysis of complex systems in many different fields (Azimi et al. 2014; Corolli et al. 2012; Azimi 2017; Okubo and Yokomori 2016; Ehrenfeucht et al. 2010, 2011). Recent applications concerned, e.g., with the efficacy of medical treatments for comorbidities and with the selection of the best environment to achieve some desired phenomena (Bowles et al. 2024; Brodo et al. 2025), led to experimenting with environments that exhibit nondeterministic and recursive behaviour. As a consequence, performing reachability and causal analysis requires the exploration of large state spaces, for which the prototype tool BioResolve (Brodo et al. 2021) struggled in terms of memory consumption and response time.

✉ Arend Rensink
arend.rensink@utwente.nl

Roberto Bruni
roberto.bruni@unipi.it

¹ CS Department, University of Pisa, Largo B. Pontecorvo, 3, 56127 Pisa, Italy

² CS Department, University of Twente, Hallenweg 19, 7522 Enschede, Netherlands

Graph Transformation (GT) (Ehrig et al. 2006; Heckel and Taentzer 2020) is a modelling technique that is widely applicable in problem domains where the objects of study have an inherent graphical structure, and the task at hand is to study their properties and evolution. Besides the graphs themselves, the core concept is that of a (transformation) rule capturing a particular change to such a graph. Rules can be used, for instance, to describe the change of a system over time, but can also be instrumental in composing and decomposing graphs and so exposing structural properties. Since RSs can be derived from Boolean network models and visualised themselves as suitable networks of reactions, it is quite natural trying to embed them within the GT framework to take advantage of well established analysis techniques.¹

Importantly, from a practical point of view, there are a number of (academic) tools supporting the use of GT. The research described here crucially relies on GROOVE (Ghamarian et al. 2012), one of the most prominent tools in this area, which was designed precisely to enable GT-based system analysis of the kind described above. The features of GROOVE that are essential for the purpose of this research are:

1. Nested (i.e., quantified) rules, which capture simultaneous changes in all neighbourhoods that satisfy certain application conditions, rather than only locally in one such neighbourhood at a time;
2. Complete exploration of the set of reachable states (under the given rules) using various strategies;
3. Model checking functionalities that can be used to validate previous findings as well to explore and support the study of new behavioural and structural properties.

The main research question that motivated our study is: *how can GT help in addressing the analysis of Reaction Systems?* To this aim, we encode a given RS as a single graph, upon which a small number of (fixed) rules can simulate the correct semantics.

The core rule describes the simultaneous firing of all **Reactions** of which no inhibitors and all reactants are present; the firing results in the presence of all products. Simultaneously, all currently present **Entity**s are removed. In GROOVE syntax, the core rule is drawn as in Fig. 1. To parse this, note that




Fig. 1 Rule for reactions firing

(in the GROOVE notation) the red structure must be absent for the rule to apply; moreover, green labels are added and blue ones deleted upon rule application. The *present* flag signals whether an **Entity** is considered to be currently present; hence, creating or deleting that flag comes down to creating or deleting the **Entity**. The \forall -nodes impose the desired quantification, causing a single application of this rule to model the firing of all enabled **Reactions**, even if there are thousands of them. Similarly for the simultaneous deletion of all **Entity**s.

Our model also supports environments that inject **Entity**s in a controlled manner. This is achieved by encoding the context specification in the initial graph and exploiting a second predefined rule, not shown here (see Fig. 7). A configuration is reachable if it can be constructed by the alternating application of both rules: first the context produces a set of stimuli (possibly upon inspection of the current state, but also nondeterministically or a combination of both), and then all enabled reactions are executed.

After a brief account of RSs (Sect. 2.1) and GROOVE (Sect. 2.2), we present the overall rationale, blueprint and key features of our approach in Sect. 4 through a simple vending machine example (Sect. 3). The subsequent experimentation in Sect. 5 is concerned with revisiting existing RS case studies to assess how GROOVE can enhance their analysis. In particular, we tackle the comorbidity treatment scenario from Bowles et al. (2024) in Sect. 5.1, the protein signalling networks analysis from Ballis et al. (2024) in Sect. 5.2, and the T cell differentiation study from Brodo et al. (2025) in Sect. 5.3. Our results are encouraging: GROOVE is capable to explore large RSs by using roughly one tenth of the time compared to BioResolve for both reachability and causal analyses. More precisely, we are able to find a trace towards unwanted patterns (if they exist) among hundreds of thousands of reachable configurations using different heuristics; then, we can also prune the trace to extract a graphical representation of the causal history of the entities. Using the built-in model checker it is also possible to confirm the results of previous studies, as well as proving new facts. A comparison with related tools is drawn in Sect. 6. Concluding remarks and future directions are discussed in Sect. 7.

¹ It should be noted that there is another, quite unrelated, way in which the concepts of Graph Transformation and Reaction Systems may be combined, namely by extending the latter from pure entities (which are analogous to graph nodes) to entities with relations between them (which can be encoded through graph edges). In Kreowski and Rozenberg (2019) this has inspired a new methodology for Graph Transformation, there called *graph surfing*.

129 2 Background

130 2.1 RSs with guarded contexts

131 First, we briefly account for the classical set theoretic definition
 132 of Reaction Systems (RSs) (Ehrenfeucht and Rozenberg
 133 2007). Then, we focus on their process algebraic version
 134 (Brodo et al. 2021) and its further extension with guarded
 135 contexts in Bowles et al. (2024) that are supported by the RS
 136 analysis framework BioResolve.²

137 *RS basics.* A Reaction System is a pair $A = (S, A)$, where S
 138 is the finite set of *entities*, and A is a finite set of *reactions* of
 139 the form $a = (R, I, P)$, with $R, I, P \subseteq S$ and $R \cap I = \emptyset$.
 140 The sets R, I, P are the sets of *reactants*, *inhibitors*, and
 141 *products*, respectively. Without loss of generality, we admit
 142 the use of empty sets as reactants, inhibitors, and products.
 143 Given the current state $W \subseteq S$, a reaction $a = (R, I, P)$ is
 144 enabled in W if all its reactants are present (i.e., $R \subseteq W$)
 145 and all its inhibitors absent (i.e., $W \cap I = \emptyset$). The *result* of
 146 the reaction a on the current state W is P if a is enabled,
 147 and \emptyset otherwise. The result of all reactions A on the current
 148 state W , is the union of the results of all reactions. The no-
 149 permanency principle of RSs dictates that entities disappear
 150 if not sustained by some reaction. Thus, the current state
 151 $W = D \cup C$ is determined by the result D of all reactions
 152 on the previous state, together with some additional entities
 153 C that can be provided by the *context* at each step.

154 *Process algebraic RSs and guarded contexts.* Inspired by
 155 Plotkin's Structural Operational Semantics approach (Plotkin
 156 2004) and process algebras such as CCS (Milner 1980), the
 157 key features of the process algebraic version of RSs are
 158 compositionality and the ability to account for a quite general
 159 notion of context (guarded, nondeterministic, recursive)
 160 using a friendly syntax. This way, we derive a Labelled Transition
 161 System (LTS) semantics for RSs by means of inductive
 162 inference rules, where LTS states are terms of an algebra,
 163 each transition defines a computation step of the RS and its
 164 label records the entities involved in that step.

165 *RS processes* are defined by the grammar below:

166 $P ::= [M]$
 167 $M ::= (R, I, P) \mid D \mid K \mid M|M$
 168 $K ::= \mathbf{0} \mid (R, I, C).K \mid K + K \mid X$

169 where R, I, P, C , and D are sets of entities (with $R \cap I = \emptyset$)
 170 and X is a context identifier drawn from a family of (possibly
 171 recursive) definitions $\Delta \triangleq \{X_j = K_j\}_{j \in J}$, called the
 172 *environment*.

173 Roughly, a RS process P embeds a *mixture* process
 174 M obtained as the parallel composition of some reactions

(R, I, P), some available entities D (if any), and some *context* process K . A context process K is either: the nil context $\mathbf{0}$ that stops the computation; the guarded context $(R, I, C).K$ that makes the entities C available to the reactions if the reactants R are present and the inhibitors I are absent, and then will behave as K at the next step; the non-deterministic choice $K_1 + K_2$ that can behave as either K_1 or K_2 ; the context identifier X that behaves as K for $X = K \in \Delta$. We write $C.K$ as a shorthand for the trivially guarded process $(\emptyset, \emptyset, C).K$ and we assume the recursive context $\text{Emp} = \emptyset.\text{Emp}$ is always defined.

We say that P and P' are structurally equivalent, written $P \equiv P'$, when they denote the same term up to the laws of Abelian monoids (unit, associativity and commutativity) for parallel composition $\cdot|.$, with \emptyset as the unit, and the laws of idempotent Abelian monoids for choice $\cdot + \cdot$, with $\mathbf{0}$ as the unit. We also assume $D_1|D_2 \equiv D_1 \cup D_2$ for any $D_1, D_2 \subseteq S$. Indexed sums and parallel compositions are denoted, respectively, by $\sum_{i \in I} K_i$ and $\prod_{i \in I} M_i$.

The SOS semantics of RS processes is defined by the SOS rules in Fig. 2. A transition label ℓ , written $\langle\langle D \triangleright R', I', C \rangle\rangle \triangleright R, I, P$, records: the available entities D ; the entities C provided by the guarded contexts, assuming all entities in R' are present and those in I' are absent; the set R of entities whose presence enables or disables some reactions; the set I of entities whose absence enables or disables some reactions; and the set P of reaction products. The rules guarantee that, whenever $P \xrightarrow{\langle\langle D \triangleright R', I', C \rangle\rangle \triangleright R, I, P} P'$, it holds that (R', I', C) is enabled in D and that (R, I, P) is enabled in $W \triangleq (D \cup C)$.

The rule (*Ent*) records the set of current entities D . By rule (*Cxt*), a guarded context process $(R, I, C).K$ makes available the entities in C if the reactants R are present in the current state and the inhibitors I are absent, and then reduces to K . Rules (*Suml*) and (*Sumr*) select a move of either the left or the right context, resp., discarding the other process. By rule (*Rec*), a context identifier X behaves according to its defining process K .

The rule (*Pro*) assumes the reaction (R, I, P) is enabled: it records its reactants, inhibitors, and products in the label, and leaves the reaction available at the next step, together with its products P . The rule (*Inh*) records in the label the reasons why the reaction (R, I, P) should not be executed: possibly some inhibiting entities ($J \subseteq I$) are present or some reactants ($Q \subseteq R$) are missing, with $J \cup Q \neq \emptyset$, as at least one cause is needed.

The rule (*Par*) puts two processes in parallel by pooling their labels and joining all labels components. We write $\ell_1 \cup \ell_2$ for the component-wise union of labels, while the sanity check $\ell_1 \frown \ell_2$ is required to guarantee that labels of reactants and inhibitors are consistent (see definitions in Fig. 2).

² <https://www.di.unipi.it/~bruni/LTSRS/>.

$$\begin{array}{c}
\frac{}{D \xrightarrow{\langle\langle D \triangleright \emptyset, \emptyset, \emptyset \rangle\rangle \triangleright \emptyset, \emptyset, \emptyset} \emptyset} \text{(Ent)} \\
\frac{}{(R, I, C).K \xrightarrow{\langle\langle \emptyset \triangleright R, I, C \rangle\rangle \triangleright \emptyset, \emptyset, \emptyset} K} \text{(Cxt)} \\
\\
\frac{K_1 \xrightarrow{\ell} K'_1}{K_1 + K_2 \xrightarrow{\ell} K'_1} \text{(Suml)} \quad \frac{K_2 \xrightarrow{\ell} K'_2}{K_1 + K_2 \xrightarrow{\ell} K'_2} \text{(Sumr)} \quad \frac{X = K \in \Delta \quad K \xrightarrow{\ell} K'}{X \xrightarrow{\ell} K'} \text{(Rec)} \\
\\
\frac{}{(R, I, P) \xrightarrow{\langle\langle \emptyset \triangleright \emptyset, \emptyset, \emptyset \rangle\rangle \triangleright R, I, P} (R, I, P)|P} \text{(Pro)} \quad \frac{J \subseteq I \quad Q \subseteq R \quad J \cup Q \neq \emptyset}{(R, I, P) \xrightarrow{\langle\langle \emptyset \triangleright \emptyset, \emptyset, \emptyset \rangle\rangle \triangleright J, Q, \emptyset} (R, I, P)} \text{(Inh)} \\
\\
\frac{M_1 \xrightarrow{\ell_1} M'_1 \quad M_2 \xrightarrow{\ell_2} M'_2 \quad \ell_1 \frown \ell_2}{M_1 + M_2 \xrightarrow{\ell_1 \cup \ell_2} M'_1 + M'_2} \text{(Par)} \quad \frac{M \xrightarrow{\langle\langle D \triangleright R', I', C \rangle\rangle \triangleright R, I, P} M' \quad R' \subseteq D \quad R \subseteq D \cup C}{[M] \xrightarrow{\langle\langle D \triangleright R', I', C \rangle\rangle \triangleright R, I, P} [M']} \text{(Sys)}
\end{array}$$

where $\ell_1 \frown \ell_2$ and $\ell_1 \cup \ell_2$ are defined as follows:

$$\begin{aligned}
\langle\langle D_1 \triangleright R'_1, I'_1, C_1 \rangle\rangle \triangleright R_1, I_1, P_1 &\frown \langle\langle D_2 \triangleright R'_2, I'_2, C_2 \rangle\rangle \triangleright R_2, I_2, P_2 \\
&\triangleq (\bigcup_{i=1,2} D_i \cup R'_i) \cap (I'_1 \cup I'_2) = \emptyset \wedge (\bigcup_{i=1,2} D_i \cup C_i \cup R_i) \cap (I_1 \cup I_2) = \emptyset \\
\langle\langle D_1 \triangleright R'_1, I'_1, C_1 \rangle\rangle \triangleright R_1, I_1, P_1 \cup \langle\langle D_2 \triangleright R'_2, I'_2, C_2 \rangle\rangle \triangleright R_2, I_2, P_2 &\triangleq \langle\langle D_1 \cup D_2 \triangleright R'_1 \cup R'_2, I'_1 \cup I'_2, C_1 \cup C_2 \rangle\rangle \triangleright R_1 \cup R_2, I_1 \cup I_2, P_1 \cup P_2
\end{aligned}$$

Fig. 2 SOS semantics of the RS processes.

Finally, the rule (Sys) checks that all the needed reactants are available in the system. Checking the absence of inhibitors is not necessary, thanks to the sanity check in rule (Par). Note that, while the enabling of $(R, I, C).K$ requires the presence of reactants R and the absence of inhibitors I w.r.t. the set of current entities D , in the case of reactions (R, I, P) , the check is performed w.r.t. the current state $W = D \cup C$. More importantly, the products C are made available immediately from the context, not at the next step. It is worthy to mention that a conditional prefixed process that is not enabled behaves as the \emptyset process.

Remark 1 It is worth noting that any RS process can be written in the form $[Rs \mid Ks \mid D]$ where $Rs = \prod_i (R_i, I_i, P_i)$ is the parallel composition of all reactions in the system, $Ks = \prod_j K_j$ is the parallel composition of all contexts and D is the set of currently available entities. Moreover, it can be proved that a generic transition has the shape: $[Rs \mid Ks \mid D] \xrightarrow{\langle\langle D \triangleright R', I', C \rangle\rangle \triangleright R, I, P} [Rs \mid Ks' \mid P]$, i.e., reactions are always preserved by transitions, the first component D of the transition label is just the set of available entities in the source state and the new result set in the target state is just the product set P observed in the label of the transition. The choices in Ks are taken by considering the set of available entities D (in the case of guarded prefixes) and will determine the context C appearing in the label as well as the continuation Ks' appearing in the target state. Given D and C , the product P is then uniquely determined by Rs . For brevity, we will sometimes draw the LTS, by recording only the strict amount of information in nodes and labels. Thus the above sample transition will be abbreviated as

$[Ks \mid D] \xrightarrow{C} [Ks' \mid P]$, assuming reactions Rs are known a priori.


A first concrete example of RS that exposes most features is presented in Sect. 3.

2.2 GT and GROOVE

Graph Transformation (GT, sometimes called Graph Rewriting) is a well-established rule-based formalism, the core of which is to specify precisely how graphs may evolve. Each rule embodies a particular change, which can be applied to a given graph (in the simplest form consisting of nodes and binary edges) by establishing where in that graph the pre-conditions of the rule are met, and then adding and deleting nodes and edges as prescribed.

In this paper, we use the so-called *algebraic approach* to graph transformation (see Ehrig et al. 2006 for a formal exposition and Heckel and Taentzer (2020) for applications in the context of software engineering); moreover, we rely on the particular flavour implemented in the tool GROOVE (Ghamarian et al. 2012; Rensink 2024). Some of the relevant features of the approach and the tool are highlighted below.

Graphs are *simple* and *typed*, meaning that there is at most one edge of a given type between any two nodes and that all nodes and edges are labelled through a morphism to a given (fixed) *type graph*. Edges are *directed* (going from their *source* to their *target*). Besides binary edges, nodes can also have *flags* (which are actually self-loops that act as additional, optional labels on nodes) and *attributes* (which are actually

**Fig. 3** Core type graph for reaction systems

binary edges whose target node is a data value, e.g., an integer or a string).

Rules, in their simplest form, consist of a left hand side (LHS) and right hand side (RHS). Rule applicability is established by *matching* the LHS to the graph in question, and where a match exists, removing nodes and edges that are in the LHS but not in the RHS, and vice versa, adding nodes and edges that are in the RHS but not in the LHS. In addition, however, GROOVE supports *quantified* rules, which can simultaneously be applied to multiple places in the same graph. An example is shown in Fig. 7 below.

Evolution of a graph is defined on the basis of a graph transformation system, which is a set of rules applied to a graph at hand, giving rise to a modified graph to which every rule can be applied again, and so forth. On top of this, GROOVE allows for *control programs* that can specify in what order rules may be applied. By exploring the potential evolution of a graph in all ways allowed by the control program, GROOVE constructs the *state space* of the graph transformation system, in the form of a *labelled transition system* consisting of all reachable graphs and the rule applications between them.

Analysis consists of the exploration of the state space for a given initial graph, rule system and (optional) control program. The exploration can be tuned by built-in strategies for searching and model checking.

The power of graph transformation lies in its generality: many systems naturally lend themselves to be modelled as graphs, and algebraic rules—especially quantified ones—provide a rich framework to specify their evolution. This is in fact our motivation for using it in the current paper: reaction systems can straightforwardly be interpreted as graphs. Figure 3 shows the core types for the relevant concepts of that interpretation. (The colours just support the visualisation and have no semantics of their own.)

Note especially the (abstract) supertype **Rule** with subtypes **Reaction** and **Step**: the former is the type for the elements of A in a Reaction System $A = (S, A)$, whereas

the latter is used to represent triples (R, I, P) in a context process K . The flag **fired** is used to mark **Rules** that have triggered in the most recent step. The set S is represented by nodes of type **Entity**; for a given **Rule**, the subsets R , I , and P of S are those **Entity**s to which there is an outgoing edge labelled reactant, inhibitor or product. The subtype **Forbidden** anticipates the principle, demonstrated later in this paper, of identifying undesirable entities and specifically searching for scenarios in which those are produced. The flag **present** is used to label the entities occurring in a state W . Finally, the structure of (guarded) context processes is captured by **State** entities, with next-edges to the **Steps** that can be non-deterministically chosen; the subsequent process after such a **Step** is determined by its outgoing move-edge. **Token** nodes are used to model which **States** are currently active.

3 Running example: a toy vending machine

To illustrate some basic concepts of RSs and, in the next section, of the proposed GROOVE encoding, we model a system composed of a student and a vending machine as a toy example. The vending machine accepts two different kinds of coins and can dispense either a cappuccino or an espresso when a coffee coin is inserted or a tea if a tea coin is inserted. A cappuccino is dispensed if some milk is available, otherwise espresso is produced. Assuming the powder for preparing coffee and tea are always present, the corresponding process can be written as follows:

$$\begin{aligned} \text{VM} \triangleq & (\{\text{ccoin}, \text{cpowder}\}, \{\text{nomilk}\}, \{\text{cappuccino}\}) \\ | & (\{\text{ccoin}, \text{cpowder}, \text{nomilk}\}, \emptyset, \{\text{espresso}\}) \\ | & (\{\text{tcoin}, \text{tpowder}\}, \emptyset, \{\text{tea}\}) \\ | & (\{\text{cpowder}\}, \emptyset, \{\text{cpowder}\}) \\ | & (\{\text{tpowder}\}, \emptyset, \{\text{tpowder}\}) \end{aligned}$$


A refill context process can, nondeterministically, refill the machine with milk.

$$\text{Refill} \triangleq \{\text{nomilk}\}. \text{Refill} + \emptyset. \text{Refill}$$

The student process is very simple: she takes cappuccino in the morning and tea in the afternoon, otherwise she gets angry.

$$\begin{aligned} \text{Student} \triangleq & (\{\text{am}\}, \emptyset, \{\text{ccoin}\}). \text{GetCappuccino} \\ & + (\emptyset, \{\text{am}\}, \{\text{tcoin}\}). \text{GetTea} \\ & + \{\text{idle}\}. \text{Student} \\ \text{GetCappuccino} \triangleq & (\{\text{cappuccino}\}, \emptyset, \emptyset). \text{Student} \\ & + (\{\text{espresso}\}, \emptyset, \{\text{anger}\}). \text{Student} \\ \text{GetTea} \triangleq & (\{\text{tea}\}, \emptyset, \emptyset). \text{Student} \\ & + (\emptyset, \{\text{tea}\}, \{\text{anger}\}). \text{Student} \end{aligned}$$

Fig. 4 LTS of the toy example. For brevity we use the notation introduced in Remark 1, where node labels only account for the list of (semicolon separated) current contexts and entities and transition labels carry the entities provided by the context



356 If the student is angry, she will bang the machine:

357 $\text{Anger} \triangleq (\{\text{anger}\}, \emptyset, \{\text{bang}\})$

358 Finally, two more reactions model the passage of time (morn-
359 ing vs afternoon) while the student is idle (i.e., not in the
360 process of getting beverages).

361 $\text{Day} \triangleq (\{\text{idle}\}, \{\text{am}\}, \{\text{am}\})$
| $(\{\text{am}\}, \{\text{idle}\}, \{\text{am}\})$

362 We assume that, initially, both entities `cpowder` and
363 `tpowder` are present. So the system can be coded as the
364 guarded RS process

365 $[\text{Refill} \mid \text{Student} \mid \{\text{cpowder}, \text{tpowder}\} \mid \text{VM} \mid \text{Anger} \mid \text{Day}]$.

366 The complete encoding of the above RS in BioResolve
367 syntax is reported in Fig. 19 in the Appendix. Using the
368 BioResolve directive `main_do(digraph)`, we can auto-
369 matically generate the underlying LTS as in Fig. 4: the initial
370 state is in light blue, while there is also a “bad” state in which
371 the student is banging the machine, shown in light coral. Note
372 that, as the entity `am` is initially not present, it means that the
373 initial time is in the afternoon. In drawing the transition sys-
374 tem, we have used the representational convention explained
375 in Remark 1).

376 In this particular example, we wish to analyse why `bang` is
377 produced, i.e., what are its causes. By manual inspection we
378 can recover a trace starting from the initial state and leading
379 to the “bad” state, e.g., $\xrightarrow{\text{idle}} \xrightarrow{\text{ccoin};\text{nomilk}} \xrightarrow{\text{anger}}$ (highlighted
380 in red in Fig. 4). From this trace, using the dynamic slicing
381 process described in Brodo et al. (2024a), we can reconstruct
382 that the production of `bang` was due to the prior production

383 of anger, which in turn was caused by the student getting
384 espresso instead of cappuccino, which is because there was
385 `nomilk` when a `ccoin` was inserted at `am`.

4 Encoding of RS in GT

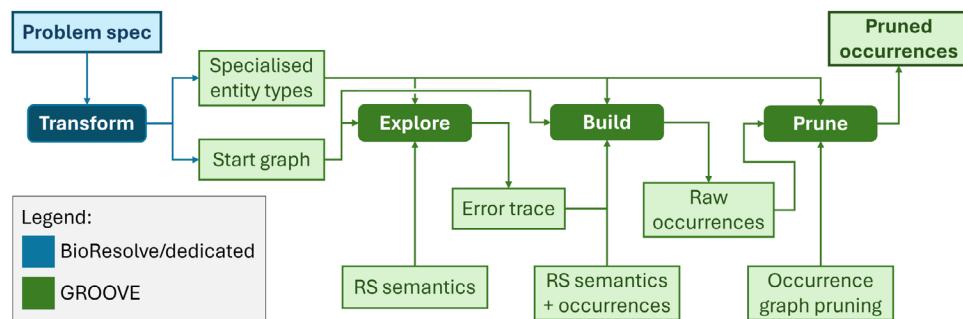
386 As an alternative to BioResolve, we investigate the use of
387 graph transformation and GROOVE to generate the underly-
388 ing LTS of a given Reaction System, on the basis of a start
389 graph obtained by transformation from the RS specification.
390 Because the start graph will typically include special **Entity**
391 subtypes, it comes together with an additional type graph
392 where those are specified. Depending on what one wants to
393 analyse, the various strengths and capabilities of GROOVE
394 then come into play.

395 For instance, one possibility is to use GROOVE’s model
396 checking capabilities to check for temporal patterns of entity
397 generation in the transition system. Another way to proceed
398 is to focus on a given trace and build its *occurrence graph*,
399 which contains all the rule occurrences and entity instances
400 present in that trace—analogous, in fact, to the way a Petri net
401 process captures a particular behaviour. If the trace in ques-
402 tion leads to a state in which a **Forbidden** entity is present
403 (such as the `bang` entity in our toy example), we can also
404 *prune* the occurrence graph, again using graph transfor-
405 mation, to keep only those rule occurrences and entity instances
406 that directly contributed to the existence of the forbidden
407 entity.

408 This gives rise to the tool chain depicted in Fig. 5, the
409 phases of which will be explained in some more detail in the
410 remainder of this section.

411 **Transform.** The first step is a text-to-model transforma-
412 tion from a problem specification in BioResolve syntax

Fig. 5 Reaction system exploration and analysis using GROOVE



414 into GROOVE syntax. This is achieved by running the
 415 `main_do(rs2gts)` directive of BioResolve, which pro-
 416 duces two artefacts: firstly, an additional type graph, com-
 417 plementary to the one shown in Fig. 3, which specifies one
 418 subtype of **Entity** for each of the entities in the problem at
 419 hand (essentially for performance reasons: relying on dedi-
 420 cated types speeds up the matching step of GROOVE); and
 421 secondly (more importantly) a start graph in which the entire
 422 BioResolve system is encoded as suggested by Fig. 3. For
 423 the example system, the additional types as well as two self-
 424 explanatory fragments of the start graph are shown in Fig. 6.
 425

We claim that this transformation is semantics-preserving;
 426 Appendix A gives a sketch of the argument. A fully formal
 427 statement and proof of semantic correspondence, however,
 428 is outside of the scope of this paper.

Explore. The dynamics of Reaction Systems is encoded as
 429 a combination of two rules, `context` and `react`, which are
 430 scheduled to fire in alternation. The rule `context` encodes the
 431 simultaneous firing of all context processes (nondeterministic-
 432 ally selecting an enabled **Step** from every **State** with a
 433 **Token**), whereas `react` encodes the (deterministic) simulta-
 434 neous firing of all enabled **Reactions**, while simultaneously
 435 erasing all **Entity**s that were not just produced. The produc-
 436 tion or erasure of an **Entity** is encoded through the creation
 437 or deletion of a *present* flag on a (persistent) **Entity** node, *not*
 438 by the creation or deletion of the node itself. In addition, to
 439 keep track of which nondeterministic choices were actually
 440 taken, the `context` rule marks the **Steps** that were selected
 441 with a *fired* flag, which is subsequently erased by the `react`
 442 rule.

Figure 7 shows the first (and most intricate) of these
 444 rules, viz. the one for the context firing. This is a quanti-
 446 fied rule, which can be read as follows: For all States with
 447 a **Token**, there is a next **Step** such that for all inhibitors
 448 there is no *present* flag whereas for all reactants there is a
 449 *present* flag; moreover, when the rule is applied, all prod-
 450 ucts of the selected **Steps** receive a *present* flag, the **Steps**
 451 themselves receive a *fired* flag, and all **Tokens** move to the
 452 successor **States**. Colour coding is used in the visual rep-
 453 resentation to distinguish the quantifier nodes \forall and \exists (both




Fig. 6 Graph representation of running example

in purple), as well as the mandatory absence (red), deletion (blue) and creation (green) of edges and flags.³

³ This colour coding is GROOVE-specific and entirely separate from the problem-specific colouring of the graph nodes in Figs. 3 and 6; in fact, to avoid confusion, the problem-specific colouring is *not* used in the rule view.




Fig. 7 Rule for context firing




Fig. 8 GROOVE LTS of the toy example

To mimic the BioResolve semantics as closely as possible, we can instruct GROOVE to regard every pair of context- and react-transitions as an atomic transaction, corresponding precisely to a transition in BioResolve (though not with the same label), and then generate the entire state space. This is achieved through a control program of the form

```
462 recipe fire() {
463   context; react;
464 }
```

where a “recipe” is the keyword for a transaction wrapping the body. With this in place, the GUI-based version of GROOVE produces the transition system displayed in Fig. 8 (which can also be exported to a range of standard formats), which is easily (visually) checked to be essentially isomorphic to the one in Fig. 4.

Alternatively, we can (for instance) ask GROOVE to search the for the first reachable state in which a **Forbidden** entity appears, using breadth-first search. (In fact, the state property for which GROOVE searches is itself determined by a *rule*, which in this case merely tests for the presence of a **Forbidden** entity). When found, the trace to the forbidden state can be saved as a control program that drives the next

stage of GROOVE exploration. In particular, using the alternating application of the context and react rules (rather than the transactional variant used for Fig. 8) this control program also records the fired-applications that tell which **Steps** have fired: this completely determines how the non-determinism in the context process has been resolved in order to arrive at the forbidden state. Here is the control program for the shortest trace to state *s*17 in our running example:

```
486 context;
487 fired("student_2");
488 fired("refill_1");
489 react;
490 context;
491 fired("student_1");
492 fired("refill_0");
493 react;
494 context;
495 fired("refill_1");
496 fired("getcappuccino_1");
497 react;
```

Here *student_2*, *student_1* and *getcappuccino_1* are the **Steps** of the Student process visualised in Fig. 6; *refill_1* and *refill_0* are the steps of the Refill process given in Sect. 3.

Build. The purpose of this phase is to build an occurrence graph that explains how **Forbidden** was produced, by collecting its (transitive) dependencies (Brodo et al. 2024b). Concretely, we record the following dependencies:

- From each non-initial **Entity** instance to the **Rule** occurrence of which it is the product;
- From each **Rule** occurrence to all its reactant **Entity** instances;
- From each **Step** occurrence to all directly preceding **Step** occurrences.

Figure 9 shows the occurrence type graph.

With respect to the slicing algorithm used in, e.g., Brodo et al. (2025), Brodo et al. (2024a), the difference is that our dependencies are based on instances, and that we explicitly include the rule occurrences and their predecessors. Each of these different kinds of dependencies is visualised both through a specific edge (product, reactant or predecessor) between the relevant instance- and occurrence-nodes, and (for the sake of a more uniform treatment during the next step of *pruning*) through an auxiliary depends-edge on the level of **TreeNode**, of which all others are subtrees.

Note that all this is restricted to *positive* dependencies. In fact, we have constructed our example so that there are no inhibitors in the **Rules** that fire in the trace above; if we would rely on an entity *milk* that inhibits the production of *espresso*, rather than on *nomilk* as a reactant, the occurrence graph for *bang* would not include *milk*; and likewise if we

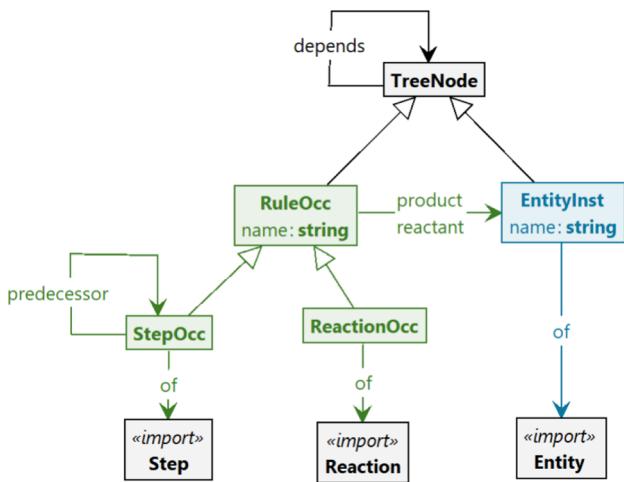


Fig. 9 Occurrence type graph

would use **cappuccino** as an inhibitor for **anger** rather than **espresso** as a reactant for it. The representation of negative dependencies is a research question in its own, and is outside the scope of this paper (although a possibility to capture negative dependencies is to focus the analysis on the Positive RS that results from the transformation defined in Brodo et al. (2024b)).

As indicated in Fig. 5, the occurrence graph is produced by another GROOVE rule system, using the same start graph but driven by the control program previously created in the explore phase, which encodes (as we have seen) the trace to the undesirable state. The occurrence graph semantics consists of rules with the same names (**react**, **context** and **fired**), but different functionality: in particular, rather than manipulating **present** flags, the **react** rule now creates **RuleOcc**- and **EntityInst**-nodes together with their dependencies. This is a non-trivial procedure that in fact itself requires several successive stages. Though the details of these stages are not of sufficient interest to include in this paper, we want to point out that breaking down a single rule (**react**, in this case) into multiple stages would seem to contradict the tenet of algebraic graph transformation that a rule embodies a single, atomic change to a graph. This contradiction is solved by relying once more on GROOVE recipes: in the occurrence graph semantics, **react** is actually not a rule but a recipe, defined as

```
555   recipe react() {
556     entities-age;
557     react-produce;
558     merge;
559   }
```

of which the three atomic steps perform the necessary book-keeping to correctly produce the occurrence graph.

Prune. The occurrence graph built by the rule system described above is too large to be useful: it contains *all* entity instances and rule occurrences produced by the trace, not just the dependencies of the undesired **Forbidden** entity. Moreover, the entire start graph is also (still) present. Therefore, in a third phase, all redundant information is pruned. This is achieved by first marking all transitive backward dependencies, and then removing all unmarked nodes. Since this is straightforward, and of no particular interest in the context of this paper, we omit the details of the GROOVE rule system for this phase. Its outcome for our running example is shown in Fig. 10 (where we have elided the auxiliary **depends**-edges).

The pruned occurrence graph visualises the causal effect chain already explained informally at the end of Sect. 3: the presence of a **ccoin**, which itself depends on **am**, combined with **cpowder** and **nomilk** causes the production of **espresso**, after which the student produces **anger** and then **bang**, which is **Forbidden**.

5 Experimentation

Here we consider three larger case studies whose RS specifications have already appeared in the recent literature. For each case study, we briefly describe its main features and then show how the methodology outlined in the previous sections can be applied for carrying out some fruitful experimentation with GROOVE.

All GROOVE experiments were carried out using GROOVE version 7.4.3 on a Dell Precision 3551 laptop with an Intel i7 CPU running at 2.6 GHz; GROOVE was run in a Java 24 JVM with 12GB of memory. No attempt was made to measure running time with precision, and repeated experiments have shown that the reported durations can deviate up to 25%. In order to facilitate replication of the experiments, we have included supplementary materials with this paper, including the required rule systems and start graphs and instructions for invoking GROOVE; see Sect. C.5.

5.1 Comorbidity treatment analysis

This case was studied in Bowles et al. (2024), where guarded contexts were introduced to handle key features of medical treatments. It concerns the risk mitigation of medication harm in the treatment of patients with comorbidities; i.e., patients with two or more long-term chronic conditions (such as diabetes, hypertension, cardiovascular diseases, chronic kidney disease, cancer, chronic obstructive pulmonary disease, among many others), who are therefore subject to follow several treatment plans simultaneously, called *clinical guidelines* (Feder et al. 1999; Woolf et al. 1999). Since clinical guidelines address a single disease, comorbidities easily lead to *polypharmacy*, where five or more medications must be




Fig. 10 Pruned occurrence graph of the running example

Features of interest. In this case study, reachability and causal analysis are key issues. Specifically, reachability is used to address questions such as *can the combination of clinical guidelines expose the patient at serious risks because of drugs interference?* Then, in the affirmative case, causal analysis can help to detect which medical decisions would be directly responsible for causing serious harm as well as to point out which alternative treatment would be available, if any. We selected this case study because the use of guarded contexts introduces new challenges for the causal analysis of RSs: while existing approaches typically focus on identifying combinations of drugs administered within the context that may have caused harm, they often fail to highlight the medical decisions that led to their administration.

Experimental set up. The RS encoding proposed in Bowles et al. (2024) relies on a formal representation of patient profiles, medical guidelines and adverse drug reactions.

For each drug d that appears in the therapies, we consider three corresponding entities get_d , stop_d and d : the first represents the prescription of d by the doctor, the second the removal of d from the current treatment and the third the intake of the drug by the patient. For handling multiple drugs of the same class c , we exploit analogous entities stop_c and c .

For medical guidelines, it takes in input the event structure modelling of therapies introduced in Bowles and Caminati (2017). Roughly, to each event e there is an identifier E_e defined as a sum of processes, one for each outgoing arc of e . If some guard is attached to the arc, then the corresponding alternative is also guarded. The prescription of a drug d is modelled by the provision of the entity get_d . Similarly, if the therapy requires stopping the drug d , the entity stop_d is produced.

The patient profile is determined by the conditions that trigger the treatment (e.g., headache, hypertension) and by the conditions that appear in the arc labels of the event structure (e.g., pregnant, asthma). We call them *features*. Correspondingly, there is one context $K_f = \{f\}.\text{Emp}$ for each feature f , and a patient profile is just a combination of some features $K_{f_1} \mid \dots \mid K_{f_n}$. Once the profile is determined by the context, it is preserved during the rest of the computation by reactions of the form $(\{f\}, \emptyset, \{f\})$, one for each feature. Accounting for all possible patient profiles within a single model can be done by considering the context $\prod_f (K_f + \text{Emp})$.

For each drug d of class c , there will be the following reactions: $(\{\text{get}_d\}, \{\text{stop}_d, \text{stop}_c\}, \{d, c\})$ modelling the intake of the drug d as for doctor prescription, and $(\{d\}, \{\text{stop}_d, \text{stop}_c\}, \{d, c\})$ modelling the prosecution of the therapy. Adverse drug reactions are provided in the form of so-called ADR tables. Each row corresponds to a set of medications M , a textual description of their side effects and

administered, increasing the risk of adverse drug reactions, or of making certain drugs less effective when combined (Hughes et al. 2013).

Analysis goals. The goal of the analysis is to explore the combination of clinical guidelines in the presence of comorbidities and for different patient profiles to detect if major risks can arise from the treatments and which profiles are exposed at severe risks. Using formal methods for risk mitigation intends to help doctors choose between alternative treatment options as well as to point out missing conditions that could be helpful to revise and update clinical guidelines.

risks when used in combination, and a severity level m (e.g., minor, moderate, major). Each row translates to a reaction $(M, \emptyset, \{m\})$.

The whole RS specification can be found in the Appendix: the list of reactions is in Fig. 20 and context processes are defined in Fig. 21.

Previous approach. The approach outlined in Bowles et al. (2024) has been used to synthesise the patient profiles that are more at risk, as a support for dynamic guideline revision: by refining guarded contexts to prevent severe effects for specific patients, we can readily check the efficacy of the changes.

GROOVE experimentation. The benefit of using GROOVE in this case study is that, besides identifying situations where a risk is found, for any risk so identified the corresponding occurrence graph of a risk can be generated, using the process outlined in Sect. 4. This provides a means for medical experts to more easily analyse root causes: for any risk that has been identified, what is the causal structure of the steps and entities leading up to it?

GROOVE can be used for full state space generation, for instance to count the number of ways a minor, moderate or major risk can arise. Some statistics can be found in Fig. 11.

The first two lines show that depth-first exploration strategy (DFS) is generally more convenient than breadth-first (BFS), so we rely on the former for the subsequent entries. At first sight, the number of major (and moderate) risks reported in Fig. 11 seems impossibly large, and evidently implies that there are patient profiles that give rise to many risks. However, this should be interpreted with care: the count refers to the total number of configurations containing a **Forbidden** (i.e., major or minor) entity, and there may very well be entities whose presence or absence does not causally contribute to that **Forbidden** entity—in other words, which would not appear in its occurrence graph. Configurations counted as separate risks may well reduce to the same causation. To analyse this further, one would have to construct (and prune) the occurrence graphs for all risk configurations, and compare them on that basis. Though this is beyond the scope of this paper, such an analysis is in principle straightforward to carry out in GROOVE—it is a matter of combining the three steps in Fig. 5 into a single rule system.

The line of Fig. 11a headed “Maj&Mod risks” reports the number of configurations at which *both* a major and a moderate entity appear during the same step; hence, these are counted as both major and moderate risks (partially explaining their high numbers).


By exploring only up to a certain depth, we can get some idea of the number of steps after which a risk typically appears, which in turn indicates the complexity of the context in which it appears. Figure 11b shows how many major risks occur after a fixed number of reaction steps, and also how much of

Measurement	Search method	Count	Time (s)
Total states	BFS	309 798	3 368
Total states	DFS	309 798	2 464
Major risks	DFS	91 113	2 447
Moderate risks	DFS	97 805	1 534
Maj&Mod risks	DFS	61 976	2 727
Minor risks	DFS	24	2 218

(a) Full exploration

Explore depth	Total risks	New risks	Total states	New states
1	0	0	769	769
2	0	0	1 345	576
3	0	0	1 873	528
4	716	716	8 133	6 260
5	9 317	8 601	42 293	34 160
6	35 899	26 582	126 012	83 719
7	73 591	37 693	227 866	101 854
8	81 691	8 100	248 420	20 554
9	88 631	6 940	276 924	28 504
10	89 733	1 102	281 854	4 930
11	90 573	840	286 054	4 200
12	91 133	560	288 854	2 800

(b) Bounded BFS exploration of major risks (tabular)



(c) Bounded BFS exploration of major risks (chart)

Fig. 11 States, risks and execution time

the total state space is involved. Note that, in this case, the total state count (reached after 12 steps) stays below that of Fig. 11b; this is because in the experiments of Fig. 11b we stop exploration at states where a risk has been found. Figure 11c shows the same data in a graphical form.

Alternatively, we can stop exploring after having found a predetermined number of risks. By setting the exploration strategy to breadth-first search, it is guaranteed that the risks found are those reached after the shortest number of steps, meaning they are the easiest to analyse visually. For instance,




Fig. 12 Occurrence graph of a major risk

Fig. 12 shows the occurrence graph of the first major risk found in this way.

The patient configuration in question is a combination of `has_fib`, `afib` and `heart_rate`; the combination of the first two leads to the prescription of `flecainide` and the second to the prescription of `diltiazem`, the combination of which should, however, be avoided. By counting the longest chain of `StepOcc`-nodes, it is confirmed that it indeed takes 4 steps to establish this risk.

We can also use GROOVE to replicate the findings of Bowles et al., (2024, Fig. 6) in terms of the relation between patient profiles and risks, using model checking. Recall that the reaction system starts by having the context produce initial entities, and in this particular case study, the first move of the context is to select a patient profile; hence the initial state has $2^9 = 512$ outgoing transitions, whose target state corresponds to the chosen profile. Moreover, the rule forbidden tests for the presence of a **Forbidden** entity in a state. Therefore, a formula of the shape $\text{AX}(\bigwedge_i f_i \rightarrow \text{EF } \text{forbidden})$, where each of the f_i specifies the presence or absence of a

patient feature, specifies whether all patient profiles with that combination of features contain a potential risk.

Concretely, in the case study at hand, Bowles et al. (2024, Fig. 6) contains necessary and sufficient criteria for patient profiles to contain major, moderate and minor risks. The part of the table pertaining to major risks is reproduced in Fig. 13. This should be read as: *precisely* in the combination of features where the green ones are present and the red ones absent, a major risk may occur.

In order to replicate these results, we can use CTL model checking. For instance, the relevant CTL property for the major risks is also shown in Fig. 13. Running the CTL model checker built into GROOVE, it reports that this is indeed satisfied, as are the corresponding characteristic properties for the moderate and minor risks. The following table reports the time taken for these checks (where the precision of our time measurement is such that the apparent difference between the model checking times is not significant):

Model generation:	2236 s
Major risk check:	67 s
Moderate risk check:	71 s
Minor risk check:	65 s

Note that the time taken for model generation is consistent with that reported in Fig. 11b.

Discussion. Compared to the prior results in Bowles et al. (2024), the advantages of using GROOVE lie in performance and flexibility:

- The time needed to analyse the entire state space is around 41 min for GROOVE; while not particularly fast, this still compares very favourably to the 5 h needed by BioResolve for LTS generation.
- Finding shortest paths to risk configurations and computing occurrence graphs is part of the core functionality of GROOVE—given, of course, suitable rule systems that encode the chosen notion of causal dependency.
- The CTL check can be used to immediately confirm the outcome of the slicing algorithm.

5.2 Protein signalling networks analysis

This case was studied in Ballis et al. (2024), where it was encoded into the Maude⁴ ecosystem (Clavel et al. 2007) to take advantage of their built-in LTL and CTL model checker facilities. It is based on a biological case study from der Heyde et al. (2014), aimed to identify the best drug treatment for three different breast cancer representative cell lines: BT474, SKBR3 and HCC1954. This is achieved by

⁴ <https://maude.cs.illinois.edu>.

Patient profiles possibly leading to a "major" adverse reaction										
	afib	has_fib	heart_rate	consensus_acei	over75	below55	diabetes	doac_int	hyper	origin
1	TRUE		TRUE	TRUE	FALSE			FALSE	TRUE	
2	TRUE		TRUE	TRUE	FALSE		TRUE	FALSE		
3	TRUE		TRUE	TRUE	TRUE	FALSE		FALSE		
4	TRUE	TRUE		TRUE				FALSE	TRUE	
5	TRUE	TRUE							TRUE	
6	TRUE	TRUE	TRUE							

```

AX((afib          & heart_rate & consensus_acei & !over75
    afib          & heart_rate & consensus_acei & !over75      & !doac_int & hyper |
    afib          & heart_rate & consensus_acei & over75 & !below55   & diabetes & !doac_int |
    afib & has_fib           & consensus_acei           & !doac_int |
    afib & has_fib
    afib & has_fib & heart_rate           & !doac_int & hyper |
    <-> EF forbidden)           & hyper |
)

```

Fig. 13 CTL encoding of the major risk profiles found in Bowles et al., (2024, Fig. 6)

797 studying the behaviour of the protein signalling networks for
 798 the HER2-positive breast cancer subtype in the presence of
 799 different combinations of monoclonal antibody drugs. In a
 800 nutshell, Maude is a high-performance reflective language
 801 and system based on equational and rewriting logic specifi-
 802 cation. The encoding of RSs is made possible by setting up
 803 a specific rewrite theory, called **ccReact**, which is expres-
 804 sive enough to capture the relevant aspects of the protein
 805 signalling networks. The analysis conducted in Ballis et al.
 806 (2024) matches previous findings, and makes it possible to
 807 readily inspect new hypotheses.

808 *Analysis goals.* The goal of the analysis is to validate or refute
 809 some behavioural hypotheses of RSs.

810 *Features of interest.* Besides reachability analysis, mostly
 811 concerned with the possibility to reach certain attractors, the
 812 distinguishing feature of this case study is the possibility to
 813 model check RSs with guarded contexts against behavioural
 814 properties written in LTL and CTL.

815 *Experimental set up.* The technique in Ballis et al. (2024)
 816 starts directly from a RS specification, which is manually
 817 coded in **ccReact** and queried using Maude state exploration
 818 techniques and built-in model checkers. Likewise, here we
 819 can just exploit the direct translation of RSs (with guarded
 820 contexts) to GROOVE presented in the previous section, i.e.,
 821 no preprocessing is necessary. The BioResolve specification
 822 is in Fig. 22 in the Appendix. The following properties have
 823 been experimented with:

- 824 1. searching for the the presence/absence of the attractor
 825 `akt` in steady states of the BT747 cell line, where the
 826 context `[k, ket]` is considered;
- 827 2. in order to observe the interactions when either erlotinib
 828 or pertuzumab are supplied, the context `[{e, egf, hrg}`.


829 `korep`] is considered and Maude reports that there exists
 830 at least one path where that treatment is successful, but
 831 not all paths avoid a steady state where `akt` is present;

- 832 3. using the context `[k, korept]`, it is shown that, regard-
 833 less the drug used, once `pdk1` is present, inevitably the
 834 steady state includes `akt`; and that `pdk1` never appears
 835 before `erbb1` is produced (which basically means that
 836 `pdk1` is a product of the activation of the `erbb1` recep-
 837 tor);
- 838 4. finally, using the context `[k, kge]`, it is shown that by
 839 permanently providing the drug erlotinib and the stimu-
 840 lus (`egf` and `hrg`), the attractor `akt` is never produced.
 841 Moreover, Maude checks that the production of `akt` can
 842 be also inhibited by providing erlotinib only when recep-
 843 tors `erbb1` and `erbb2` are active.

844 *Previous approach.* **ccReact** allowed to perform reachabil-
 845 ity analysis directly exploiting the `search` command of
 846 Maude. The formal verification of temporal formulas has
 847 been made possible by relying on a general interface to
 848 different model checkers for Maude models, called the Uni-
 849 fied Maude Model-Checking tool (`umaudemc`) Rubio et al.
 850 (2021). Some examples of verified temporal formulas are
 851 those expressing properties such as: *Does there exist at least
 852 one path where that treatment is successful? Do all paths
 853 prevent reaching a steady state in which a AKT is present?*
 854 *GROOVE experimentation.*

855 Like Maude, GROOVE has built-in model checking capa-
 856 bilities for both LTL and CTL properties; below, we show
 857 how to replicate the results of Ballis et al. (2024), for the
 858 four scenarios listed above.

859 The main challenge in replicating the results is that some
 860 of the properties to be checked are formulated in terms of
 861 *steady states* of the reaction system, which are essentially



Tests if simulation has started and the set of input entities equals the set of present entities




Fig. 14 Additional rules and condition for detecting steadiness

one-state attractors, that is, states in which the context and reactions together reproduce exactly the entities of that state again. Though GROOVE detects such a loop as a matter of course, it is a structural property of the LTS and not a state property available for model checking. In order to be able to reason about steadiness, we have to remember *input entities*, i.e., entities that were present in the source state, and compare them to *present entities*, i.e., those that have been produced in the target state. Moreover, we should not accidentally mark the start state as steady even if it has neither inputs nor present entities. Figure 14 shows the additional rules that achieve this, together with the modified recipe defined by

```
874 recipe fire() {
875   try testStart; markInput; context; react;
876 }
```

877 The resulting steady condition is given (using quantifier syntax) in Fig. 14c.

- 879 1. Given the context [k, ket], GROOVE confirms the status
880 of the following LTL properties:
 - 881 • FG(steady \rightarrow akt) is not satisfied; GROOVE
882 produces a counter-example.
 - 883 • G(erbb2 \rightarrow X(erbb2)) is satisfied.
- 884 2. Given the context [$\{e, egf, hrg\}.korep$], GROOVE
885 confirms the status of the following CTL properties:
 - 886 • EF (steady & !akt) is satisfied;
 - 887 • AF (steady & !akt) is not satisfied.
- 888 3. Given the context [$k, korept$], GROOVE confirms the
889 status of the following LTL properties:
 - 890 • G(pdk1 \rightarrow FG(steady \rightarrow akt)) is satisfied;




Fig. 15 GTS for cancer scenario 4

- erbb1 R !pdk1 is satisfied.
- 4. Given the context [k, kge], GROOVE confirms the status
of the following CTL property:
 - EG EF (steady \rightarrow !akt) is satisfied.

However, we want to point out that this property does not actually provide any useful guarantees, because the predicate *steady* (both in Ballis et al. (2024) and in our encoding explained above) only tests for *single-state* attractors. If the reaction system ends up in a multi-state loop, *steady* will never be satisfied and hence the implication *steady* \rightarrow *!akt* is *always* satisfied, regardless of whether or not *akt* holds. Indeed, the state space of this scenario, visually reproduced in Fig. 15, has such a multi-state attractor (consisting of states $s15$ and $s18$); in both of those states the predicate *akt* holds, yet the state space as a whole satisfies the CTL property above.

In replicating the results from Ballis et al. (2024), we have had to make a few adjustments. The LTL formulas reported in Ballis et al., (2024, Page 14) for scenarios 1 and 3 are actually *not* literally the ones above, but use the predicate *io-state* rather than *isSteady*. We believe that the use of *isSteady* (or, in our case *Steady*) is more informative and probably the intended version.

As a final observation, we note that the numbers of states in all these scenarios is actually quite small. We have already shown the 6-state scenario 4 in Fig. 15; the size of the others is given by the following table.

Nr.	Context	States
1	[k, ket]	4
2	[$\{e, egf, hrg\}.korep$]	10
3	[$k, korept$]	32
4	[k, kge]	6

Discussion. Compared to the prior results in Ballis et al. (2024), the advantages of using GROOVE lie in the combination of visual inspection and automatic model checking.



Fig. 16 Graphical representation of the Boolean network model of T cell differentiation from Puniya et al. (2018)

Not only were we able to confirm the findings of the original paper using exactly the same encoding of reaction system as for the previous case, but the ability to inspect and visualise the state spaces also gives additional insights, such as the observation above that steadiness as formalised there does not actually capture the intended notion of being an attractor.

5.3 T cell differentiation analysis

The paper (Brodo et al. 2025) (being the full and corrected version of Brodo et al. (2025), see also Footnote 5) exploits Reaction Systems to analyse T cell differentiation in the immune system, a widely studied biological phenomenon. The starting point for the analysis is a Boolean network model; several of those are available as a Saez-Rodriguez et al. (2007); Thakar and Albert (2010); Puniya et al. (2018), among which the one in Puniya et al. (2018) was selected. The model encompasses reactions enabling T cells to manifest various phenotypes in response to environmental stimuli, and describes a realistic regulation system that is involved in many diseases (Lafaille 1998; Hirahara and Nakayama 2016; Meng et al. 2016).

The Boolean network model is graphically represented as shown in Fig. 16, where the 9 orange nodes represent different environmental stimuli that the T cell can receive; all the other nodes represent so-called *transcription factors* and have an associated Boolean update formula that specifies when they are triggered. T cells can manifest four phenotypes, represented by the four transcription factors *tbet*, *gata3*, *rorgt* and *foxp3*, respectively. There exists experimental and computational evidence that a T cell can manifest more than one phenotype (Luckheeram et al. 2012; Puniya et al. 2018).

Analysis goals. The reachability analysis must take into account the different combinations of phenotypes that a T cell can express, called a *target* (hence, $2^4 = 16$ targets overall). For example, for the target containing the combination of transcription factors *{tbet, gata3}*, we must select an attractor that includes at least one state in which *tbet* is present, at least one state (possibly the same) in which *gata3* is present, and no state in which either *rorgt* or *foxp3* are present. The causal analysis aims to collect the

combinations of environmental stimuli that are responsible for leading to that target.

Features of interest. We have selected this case study because it shows the applicability of our method to Boolean networks models, like those available in the public database on the CellCollective platform (Helikar et al. 2012). For these models, the most relevant viewpoints are often reachability (e.g., *which phenotypes are reachable?*) and causality (*what is the effect of environmental conditions?*) analyses. Their corresponding RSs always use a special kind of nondeterministic persistent context, where at the beginning of the experiment a subset of external stimuli is chosen and then provided at each subsequent step, inevitably causing the RS to end up in a loop (called an attractor).

Experimental set up. The translation from Boolean networks to RS consists in turning every update formula into disjunctive normal form. Then, every clause of the disjunction produces a reaction in which (i) reactants are the positive atoms, (ii) inhibitors are the negated atoms and (iii) the updated variable forms a singleton product. The translation from Boolean network to BioResolve syntax is done using the directive `main_do(bn2rs)`. For the readers' convenience, all update formulas and the resulting reactions are reported, respectively, in Fig. 22 and in Fig. 24 in the Appendix. For example, the update formula for *IL12R* is

$$\begin{aligned} & (\text{IL12\&NFAT}) \mid (\text{STAT4\&\neg GATA3}) \\ & \mid \text{Tbet} \mid (\text{TCR\&\neg GATA3}) \end{aligned}$$

which yields the four reactions⁵

$$\begin{array}{ll} (\{\text{il12, nfat}\}, \emptyset, \{\text{il12r}\}) & (\{\text{stat4}, \{\text{gata3}\}, \{\text{il12r}\}\}) \\ (\{\text{tbet}\}, \emptyset, \{\text{il12r}\}) & (\{\text{tcr}\}, \{\text{gata3}\}, \{\text{il12r}\}) \end{array}$$

The RS context can choose any combination of environmental stimuli that will then persist, i.e., for each possible stimulus s we define the context processes $X_s \triangleq \{s\}.X_s$ and then take the context $\prod_s (X_s + \text{Emp})$. The resulting LTS has an initial branching into 2^9 different states, because there are 9 possible stimuli to be considered. Subsequently, each of the 2^9 states originates a deterministic computation, leading to some attractors.

Previous approach. The paper (Brodo et al. 2025) presents a toolchain (BioResolve, SWI-Prolog, Python and Python-to-Prolog binding facilitated by the `swiplserver` Python

⁵ The specification analysed in Brodo et al. (2025), on which this paper is based, repairs a minor typo in the original conference version (Brodo et al. 2025). Specifically, the product set of the reaction *(stat4, gata3, il12r)* was mistakenly written as *il2r*, omitting the digit 1. Unfortunately, since *il2r* was also a valid entity, the error was difficult to detect. Though the mistake mildly affected the original results, it turns out that for the analysis reported in this paper there is no difference at all.: 999, 1000, 1001

package) to study the Boolean network model. Roughly, after translating the Boolean network model to RS specifications the whole LTS is constructed according to any combination of persistent stimuli that can be provided by the context. BioResolve returns the LTS as a graph in dot format, which is then loaded by a Python script. Then, attractors related with a target of interest are identified by looking for cycles in the LTS, and a slicing algorithm performs some form of causal analysis, to simplify each computation trace by preserving only the relevant causes of those target entities. The generation of the LTS is often the bottleneck of the approach, both in terms of time (Prolog performance), but also in terms of space, because BioResolve can require to allocate a large stack limit size to succeed.

GROOVE experimentation. The capabilities of GROOVE called upon for this case study are very similar to those in Sect. 5.1; the main difference lies in the specific interest in attractors. Indeed, in contrast to the situation for comorbidities, here after the initial selection of a profile, the context does not cause any more nondeterminism; hence every profile eventually ends up in such an attractor.

A trace ending in a loop, sometimes called a “lollipop”, is in fact precisely what LTL properties are checked over; hence an LTL property violation takes the form of a lollipop. This means that we can find attractors with specific properties by formulating their non-existence in LTL, and then finding a counter-example through model checking. For instance, the following formulas deny the reachability of an attractor in which `tbet` and `gata3` are expressed:

- $\neg G(F \text{ } gata3 \wedge F \text{ } tbet)$ (separate expression)
- $\neg GF(gata3 \wedge tbet)$ (simultaneous expression)

Using the start graph derived from BioResolve using the process outlined in Fig. 5, both of these formulas yield counterexamples, meaning that `tbet` and `gata3` *can* in fact be (recurrently) expressed simultaneously. Using a variation of the process outlined in Sect. 4, we can once more visualise a trace leading to such a recurrent state. The variation lies in the fact that, this time, we do not want to show the causal history of a *single* forbidden entity, but rather of the combination of two distinct entities. Fortunately, this is just a matter of creating another rule, `gata3-tbet`, which applies precisely when `gata3` & `tbet` holds. On this basis we can go through the steps outlined in Fig. 5, resulting in the occurrence graph displayed in Fig. 17.

This complements the observation embodied in Brodo et al., (2025, Fig. 7) that for the combination of `tbet` and `gata3`, the context has to provide the stimuli `ifnge` (produced here by the **StepOcc** named `x81_0`), `tcr` (repeatedly produced by **StepOccs** named `x91_0`) and `i14e` (also repeatedly produced, by **StepOccs** named `x51_0`). In more

detail, we see that `tbet` derives, in a linear sequence of four **ReactionOccs**, from `ifnge` and `tcr`, whereas `gata3` derives, in 3 successive combinations of simultaneous **ReactionOccs**, from `tcr` and `i14e`. Moreover, the genes produced along the way are precisely the ones reported in Brodo et al., (2025, Fig. 8(a)), using the slicing algorithm of that paper, as being relevant for the expression of `tbet` and `gata3`.

Besides the production of such occurrence graphs for specific cases, GROOVE can also be used once more to directly confirm the findings of Brodo et al., (2025, Figs. 7 and 8), by expressing them as CTL formulas similar to the one reported in Fig. 13. In this context, it is relevant to report that, in contrast to BioResolve, where (as reported above) time and space performance were a bottleneck in the analysis of this model, GROOVE can fully explore the state space in approximately 3 s.

Discussion Compared to the results in Brodo et al. (2025), the advantages of GROOVE are threefold (reiterating the observations made for the previous two cases):

- LTL model checking allows to express, in a flexible manner, the scenarios one wants to investigate;
- The occurrence graph visualisation offers analysis possibilities beyond the outcome of the slicing algorithm;
- The performance of GROOVE is an order of magnitude better than that of BioResolve

6 Comparison with existing tools

The approach presented in this paper can provide several advantages over existing tools in the literature.⁶


- `brsim`⁷ (Basic Reaction System Simulator, written in Haskell and distributed under the terms of GNU GPLv3 license) Azimi et al. (2015) was the first RS simulator to be made publicly available. Given the reactions of the RS and a context sequence, `brsim` is capable of computing the resulting sequence and generating additional annotations for each computation step, such as the enabled reactions. Alternatively, `brsim` can be executed in an interactive mode, allowing the user to manually provide the context to be used at each step.
- WebRSim⁸ is a basic RS simulator that makes all functionalities of `brsim` available through a friendly web interface (Ivanov et al. 2018);

⁶ See, e.g., the list of Reaction Systems Computer Environments at <https://www.reactionsystems.org/about-reaction-systems>.

⁷ Available at <https://github.com/scolobb/brsim/>.

⁸ Available at <https://github.com/scolobb/brsim>.

Fig. 17 Occurrence graph for the simultaneous expression of gata3 and tbet



- 1094 • HERESY⁹ is a Highly Efficient REaction SYstem GPU-
 1095 based simulator, developed using CUDA (Nobile et al.
 1096 2017). It features a user-friendly GUI and is designed to
 1097 exploit the high degree of parallelism offered by modern
 1098 GPUs to handle very large-scale RSs simulations.
 1099 • cl-rs¹⁰ is an optimised Common Lisp simulator for
 1100 RSs (Ferretti et al. 2020) that can exhibit performances
 1101 comparable with the GPU-based simulator HERESY. This
 1102 is achieved by discarding all reactions that cannot pro-

duce effects and by encoding RS evolution in terms of matrix–vector multiplications and vector additions.

- 1103 • BioResolve¹¹ is a Prolog interpreter for Reaction Sys-
 1104 tem analysis, first proposed in Brodo et al. (2021) and
 1105 later extended in a series of papers to deal with enhanced
 1106 features, like delays, duration, monitoring, slicing and
 1107 guarded contexts (Brodo et al. 2023, 2024a; Bowles et al.
 1108 2024). Many capabilities of BioResolve have been dis-
 1109 cussed at length in the previous sections.

⁹ Available at <https://github.com/aresio/HERESY>.

¹⁰ Available at <https://github.com/mnzuca/cl-rs>.

¹¹ Available at <https://www.di.unipi.it/~bruni/LTSRS/>.

- 1112 • **ccReact**¹² is an interacting language for Reaction Systems
1113 based on Maude 3.2.1 (Ballis et al. 2024), whose
1114 key features have been illustrated in Sect. 5.2.
- 1115 • **ReactlCS**¹³ is a Reaction Systems Verification Toolkit
1116 that consists of two main modules for model check-
1117 ing temporal properties expressed in logical languages
1118 tailored to Reaction Systems: one that exploits binary
1119 decision diagrams (BDD) and bounded model checking;
1120 the other that translates verification problems in rsLTL
1121 (Meski et al. 2015) into satisfiability modulo theories
1122 (SMT) (Meski et al. 2024).

1123 *Modelling capabilities.* The GROOVE-based method sup-
1124 ports a rich and expressive encoding of RSs, including the
1125 most recent features such as the handling of *guarded*, *recurs-
1126 ives*, and *nondeterministic contexts*. Among the other tools,
1127 such features are only supported by BioResolve, which,
1128 however, relies on a Prolog back-end that limits scalability
1129 and requires external scripting for improving the perfor-
1130 mance of many analyses whenever large state generation and
1131 exploration is necessitated. Tools such as HERESY, WebR-
1132 Sim, and cl-rs provide lightweight RS simulators but
1133 are limited to basic semantics, lacking support for more
1134 advanced interactions with the context or advanced veri-
1135 fication features. **ccReact** supports temporal logic model
1136 checking (LTL/CTL), but not recursive contexts. Moreover,
1137 the encoding of RSs in **ccReact** is manual and less suited to
1138 visual inspection or dynamic causal analysis. In addition to
1139 (bounded) model checking of custom temporal logics specif-
1140 ically designed for Reaction Systems, ReactlCS also supports
1141 context automata (a slightly less general notion of contexts
1142 than the one considered here), reactions with concentration
1143 levels, parameter and reaction synthesis, as well as parameter
1144 optimisation.

1145 *Performance and scalability.* The ability of GROOVE to
1146 explore large state spaces efficiently is central to our method.
1147 Through configurable exploration strategies and a rule-based
1148 control mechanism, GROOVE handles complex RS instances
1149 that involve thousands of reachable configurations. Our
1150 experiments demonstrate a substantial improvement in anal-
1151 ysis time compared to BioResolve, often reducing execution
1152 time by an order of magnitude. Furthermore, the performance
1153 of BioResolve is strongly influenced by the nature of Prolog
1154 evaluation strategies, which can lead to excessive memory
1155 and time consumption in large case studies. With respect to
1156 **ccReact**, the paper (Ballis et al. 2024) on which we based our
1157 experiments does not provide information on performance,
1158 and indeed the system studied there is so small that no use-
1159 ful comparison could be made on that basis. Since Maude,

1160 underlying **ccReact**, is a long-standing and mature tool, it
1161 would be interesting to investigate this in more detail; how-
1162 ever, we leave this to future work. In contrast, other tools
1163 either consider linear executions only and do not scale to
1164 large models or lack optimisation strategies necessary for
1165 handling non-trivial state spaces.

1166 Contrary to GROOVE, which is a general purpose, explicit-
1167 state tool, not optimised for Reaction Systems, ReactlCS can
1168 take advantage of symbolic representations to abstract much
1169 larger state-spaces. For instance, for the mutual exclusion
1170 protocol reported in Meski et al. (2015), Nobile et al. (2017),
1171 ReactlCS can cope with 37 to 54 processes (depending on the
1172 precise analysis) in a matter of hours; in contrast, in compara-
1173 ble time GROOVE can exhaustively generate the state space of
1174 up to 13 processes. Even though the model checking queries
1175 in ReactlCS probably do not require full state space explo-
1176 ration, and hence the comparison is skewed, the dedicated
1177 methods of ReactlCS clearly pay off.

1178 *Causal analysis and verification.* A key distinguishing fea-
1179 ture of our approach is the ability to perform graph-based
1180 *causal slicing*. By automatically generating and pruning
1181 *occurrence graphs*, GROOVE provides detailed and visual
1182 explanations of how specific states, such as those involving
1183 undesirable or forbidden entities, are reached. This form of
1184 causal reasoning is not available in high-performance tools
1185 such as HERESY, WebRSim, cl-rs, or ReactlCS and is
1186 only partially addressed in **ccReact**, where the focus is pri-
1187 marily on reachability and temporal properties. GROOVE’s
1188 integrated support for CTL and LTL model checking further
1189 extends its applicability to behavioural verification, enabling
1190 the specification and validation of complex temporal prop-
1191 erties.

1192 *Summary.* In conclusion, the combination of expressive mod-
1193 elling, efficient state space exploration, and integrated causal
1194 analysis makes GROOVE a powerful and versatile full-fledged
1195 platform for the study of Reaction Systems. It not only gen-
1196 eralises and extends existing tools, but also opens the door
1197 to new forms of analysis that were previously impractical or
1198 unsupported.

7 Conclusion and future work

1199 In this work, we have demonstrated how Reaction Systems
1200 can be effectively encoded and analysed within the GROOVE
1201 framework, so to reuse the expressiveness and efficiency
1202 of graph transformation techniques. By exploiting quanti-
1203 fied rules, the encoding consists of a direct translation from
1204 RS specification to a typed graph, which is made auto-
1205 matic in BioResolve. Then, GROOVE enables both exhaustive
1206 state space exploration, the extraction of causal information

¹² Available at <https://depot.lipn.univ-paris13.fr/olare/reaction-systems-maude>.

¹³ Available at <https://github.com/arturmeski/reactics/>.

1208 through occurrence graphs and property-based verification
 1209 based on model checking.

1210 We have used GROOVE to revisit several case studies from
 1211 the literature and our experimental results, although preliminary,
 1212 are promising: GROOVE not only supports complex
 1213 RS features such as guarded, nondeterministic and recursive
 1214 contexts but also significantly improves performance
 1215 and flexibility compared to existing RS tools. Moreover, the
 1216 use of GROOVE’s recipes and model-checking capabilities
 1217 opens the door to sophisticated analyses that were previously
 1218 impractical.

1219 A number of interesting avenues for future research
 1220 remain open, among which we mention the possibility to
 1221 extend the methodology to support quantitative RS variants
 1222 with durations or weights (Brodo et al. 2023); investigate
 1223 alternative notions of causality and their representation
 1224 in graph-based semantics, like dependencies drawn from
 1225 inhibitors rather than reactants¹⁴; apply the GROOVE toolset
 1226 to further biological case studies, like those available in the
 1227 CellCollective public repository (Helikar et al. 2012), possibly
 1228 exploiting an automated pipeline for analysis.

1229 Overall, the results show that graph transformation, and
 1230 GROOVE in particular, provide a robust and scalable foundation
 1231 for the specification, execution, and analysis of Reaction
 1232 Systems.

1233 One option that we have ignored throughout the paper
 1234 deserves a brief mention here. The slicing algorithms
 1235 reported in (Bowles et al. 2024; Brodo et al. 2025) on the
 1236 basis of BioResolve are actually implemented as stand-alone
 1237 scripts that operate on a .dot-formatted LTS. Given that
 1238 GROOVE can also produce LTSs as .dot files, an alternative
 1239 way to benefit from its superior performance might be
 1240 merely to replace BioResolve in the tool chain used previously.
 1241 For this to be possible, the LTS labelling information
 1242 produced by GROOVE should conform to the requirements
 1243 of the slicing algorithm, following the principles outlined in
 1244 Remark 1. Though that is currently not the case (compare
 1245 GROOVE’s Fig. 8 to BioResolve’s Fig. 4), we believe that
 1246 this requires only a minor adjustment of the rule system.

1247 From the perspective of GROOVE development, carrying out
 1248 the experiments described in this paper has led to many small
 1249 improvements as well as inspiration for new features. For
 1250 instance, the various strategies for “ordinary” state space
 1251 exploration (DFS- or BFS based, bounded or not, conditional
 1252 or not) on the one hand and the model checking capabilities
 1253 on the other are not well-integrated. Queries such as
 1254 “count the number of states of max depth n where a given
 1255 CTL property holds” or “find all prefixes of length n of paths
 1256 satisfying a given LTL property,” which would enhance the

1257 capabilities for analysing graph-based models such as the
 1258 ones studied here, currently cannot be posed in a straightforward
 1259 manner. Another useful extension would be the ability
 1260 to automatically chain different transformations, such as the
 1261 three steps of explore–build–prune in Fig. 5, which currently
 1262 have to be invoked separately. We plan to investigate these
 1263 extensions in the future.

Supplementary information

1264 For replication of the GROOVE experiments, we provide an
 1265 archive with supplementary material, described in Sect. C.5.

Appendix A Semantic correspondence

1266 In this paper, we claim, but do not prove, that the GROOVE
 1267 encoding correctly captures the operational Reaction Systems
 1268 semantics recalled in Sect. 2.1. A complete formal proof
 1269 is outside the scope of this paper, which focusses on exper-
 1270 imental results; however, in this section we provide a sketch
 1271 of such a proof.

1272 As described in Sect. 4, each GROOVE state is a graph
 1273 consisting of a fixed part (which is the same in every state) and
 1274 a variable part. The fixed part contains an encoding of (i) the
 1275 reaction system A itself, with the entities S as **Entity**-typed
 1276 nodes, and (ii) the automata of the context processes, with
 1277 individual processes corresponding to **Token**-typed nodes
 1278 and states as **State**-typed nodes.


1279 The variable part consisting of *present* flags on **Entity**
 1280 nodes and a current-edge from every **Token** to a **State**. As
 1281 observed in Remark 1, any RS process $[M]$ can be written
 1282 in the form $[Rs \mid Ks \mid D]$ where: (i) $Rs = \prod_i (R_i, I_i, P_i)$ is
 1283 the parallel composition of all reactions in the system, and
 1284 never changes along the computation; (ii) $Ks = \prod_j K_j$ is the
 1285 parallel composition of all contexts, and (iii) D is the set of
 1286 currently present entities.

1287 A GROOVE state G is equivalent to an RS process
 1288 $[Rs \mid Ks \mid D]$ if: (i) an entity e is present in D if and only
 1289 if its **Entity** node is flagged as *present* in G , and (ii) there
 1290 is a one-to-one correspondence of the available contexts K_j
 1291 in the RS process and the **Token**-nodes in G , such that the
 1292 current state of that **Token** is the start state of the automaton
 1293 for K_j .

1294 We claim that this equivalence establishes a bisimulation
 1295 between the GROOVE and RS state spaces. (It is not an iso-
 1296 morphism, because the RS semantics for K “consumes” the
 1297 process (rules $(Cxt)-(Rec)$ in Fig. 2) whereas G keeps the
 1298 automaton intact, merely moving the current pointer.)

1299 To prove this claim, we have to show correspondence of
 1300 the transitions. Concretely, to any transition carrying the label
 1301 $\langle \langle D \triangleright R', I', C \rangle \triangleright R, I, P \rangle$, there corresponds a single

¹⁴ In this respect, we could, e.g., exploit the semantic-preserving transformation from RSs to Positive RSs proposed in Brodo et al. (2024b).

**Fig. 18** Rule for reaction firing

application of the recipe fire, which in turn consists of applications of rules context (see Fig. 7) followed by react (see Fig. 18).

Both of these rules are universally quantified. Roughly speaking, referring to Fig. 2, context simultaneously encodes all sub-transitions of the form

$$\begin{array}{c} D \xrightarrow{\langle\langle D \triangleright \emptyset, \emptyset, \emptyset \rangle\rangle \triangleright \emptyset, \emptyset, \emptyset} D' \quad (\text{rule (Ent)}) \\ K \xrightarrow{\langle\langle \emptyset \triangleright R', I', C \rangle\rangle \triangleright \emptyset, \emptyset, \emptyset} K' \quad (\text{rules (Cxt)-(Rec)}) \end{array}$$

as well as their composition and filtering (rules (Par)–(Sys)). react in turns simultaneously encodes all sub-transitions of the form

$$(R, I, P) \xrightarrow{\langle\langle \emptyset \triangleright \emptyset, \emptyset, \emptyset \rangle\rangle \triangleright R', I', P'} M' \quad (\text{rules (Pro)-(Inh)})$$

together with their composition with the context-transitions (rule (Par)) and filtering (rule (Sys)). (The *fired*-flags on the Step-nodes, added by context and removed again by react, play no role in this correspondence: for the purpose of the equivalence of the GROOVE and RS semantics, they might be omitted entirely.)

Appendix B Mutual exclusion

Table 1 shows the GROOVE performance in generating the full state space of the mutual exclusion example of Meski et al. (2015), Nobile et al. (2017), discussed in Sect. 6.

Table 1 State space generation for mutual exclusion

Case #	States #	Time Ratio	Time s	Memory MB	Memory Ratio
2	11		0.2		1
3	27	2.5	0.4	1.9	2 2.2
4	63	2.3	0.7	2.0	4 2.1
5	143	2.3	2.0	2.8	13 3.0
6	319	2.2	5.6	2.7	34 2.6
7	703	2.2	17.5	3.2	96 2.8
8	1535	2.2	56.6	3.2	262 2.7
9	3327	2.2	175.0	3.1	700 2.7
10	7167	2.2	595.3	3.4	1,929 2.8
11	15359	2.1	1,890.3	3.2	4,764 2.5
12	32767	2.1	5,146.1	2.7	8,588 1.8
13	69631	2.1	15,647.3	3.0	13,640 1.6

```

myentities([cpowder,tpowder]). % initial set D0

myreactions([
    react([idle],[am],[am]), % list of reactions
    react([am],[idle],[am]),
    react([ccoin,cpowder],[nomilk],[cappuccino]),
    react([ccoin,cpowder,nomilk],[],[espresso]),
    react([tcoin,tpowder],[],[tea]),
    react([cpowder],[],[tpowder]),
    react([tpowder],[],[tpowder]),
    react([anger],[],[bang]) ]).

mycontext("[refill,student]"). % context processes

myenvironment("[ % context definitions
    refill = ({nomilk}.refill
        + {} . refill),
    student = (?{}, {am} ; {tcoin})? . gettea
        + ?{am}, {}, {ccoin}? . getcappuccino
        + {idle} . student,
    gettea = (?{tea}, {}, {} ? . student
        + ?{}, {tea} ; {anger} ? . student),
    getcappuccino = (?{cappuccino}, {}, {} ? . student
        + ?{espresso}, {}, {anger} ? . student) ]").

```

Fig. 19 BioResolve implementation of the vending machine RS from Sect. 3. The question marks ? are used to delimit guarded prefixes in context processes

Appendix C Auxiliary material

C.1 Auxiliary material for the toy running example

The BioResolve specification for the toy running example about the interaction between the student and the vending machine is reported in Fig. 19. The corresponding RS has been described in Sect. 3 and it has been used to illustrate some key features of the GROOVE encoding in Sect. 4.

```

Feats  $\triangleq$  ((hyper),  $\emptyset$ , {hyper}) | ((afib),  $\emptyset$ , {afib}) | ((has_fib),  $\emptyset$ , {has_fib}) | ((heart_rate),  $\emptyset$ , {heart_rate}) | ((consensus_acei),  $\emptyset$ , {consensus_acei})
| ((over75),  $\emptyset$ , {over75}) | ((below55),  $\emptyset$ , {below55}) | ((diabete),  $\emptyset$ , {diabete}) | ((origin),  $\emptyset$ , {origin})
| ((doac_int),  $\emptyset$ , {doac_int}) | ((hyper),  $\emptyset$ , {diseases}) | ((diabete),  $\emptyset$ , {diseases})

Drugs  $\triangleq$  ((get_diltiazem), {stop_cbb}, {diltiazem, cbb}) | ((diltiazem), {stop_cbb}, {diltiazem, cbb}) | ((get_verapamil), {stop_cbb}, {verapamil, cbb})
| ((verapamil), {stop_cbb}, {verapamil, cbb}) | ((diltiazem, verapamil), {stop_cbb}, {alert_dup}) | ((get_propranolol), {stop_nsbb}, {propranolol, nsbb})
| ((propranolol), {stop_nsbb}, {propranolol, nsbb}) | ((get_carvediol), {stop_nsbb}, {carvediol, nsbb}) | ((carvediol), {stop_nsbb}, {carvediol, nsbb})
| ((propranolol, carvediol), {stop_nsbb}, {alert_dup}) | ((get_bisoprolol), {stop_sbb}, {bisoprolol, sbb}) | ((bisoprolol), {stop_sbb}, {bisoprolol, sbb})
| ((get_atenolol), {stop_sbb}, {atenolol, sbb}) | ((atenolol), {stop_sbb}, {atenolol, sbb}) | ((bisoprolol, atenolol), {stop_sbb}, {alert_dup})
| ((get_flecainide), {stop_flec}, {flecainide}) | ((flecainide), {stop_flec}, {flecainide}) | ((get_warfarin), {stop_warf}, {warfarin})
| ((warfarin), {stop_warf}, {warfarin}) | ((get_apixaban), {stop_doac}, {apixaban, doac}) | ((apixaban), {stop_doac}, {apixaban, doac})
| ((get_dabigatran), {stop_doac}, {dabigatran, doac}) | ((dabigatran), {stop_doac}, {dabigatran, doac}) | ((apixaban, dabigatran), {stop_doac}, {alert_dup})
| ((get_vkant), {stop_vkant}, {vkant}) | ((vkant), {stop_vkant}, {vkant}) | ((get_benzepril), {stop_acei}, {benzepril, acei})
| ((benazepril), {stop_acei}, {benazepril, acei}) | ((get_captopril), {stop_acei}, {captopril, acei}) | ((captopril), {stop_acei}, {captopril, acei})
| ((benazepril, captopril), {stop_acei}, {alert_dup}) | ((get_olmesortan), {stop_arb}, {olmesortan, arb}) | ((olmesortan), {stop_arb}, {olmesortan, arb})
| ((get_irbesartan), {stop_arb}, {irbesartan, arb}) | ((irbesartan), {stop_arb}, {irbesartan, arb}) | ((olmesortan, irbesartan), {stop_arb}, {alert_dup})
| ((get_indapamide), {stop_td}, {indapamide, td}) | ((indapamide), {stop_td}, {indapamide, td}) | ((get_chlorothiazide), {stop_td}, {chlorothiazide, td})
| ((chlorothiazide), {stop_td}, {chlorothiazide, td}) | ((indapamide, chlorothiazide), {stop_td}, {alert_dup}) | ((doac), {doac_ok}, {doac_fail}, {doac_test})
| ((doac_ok), {doac_fail}, {doac_ok}) | ((doac_fail), {doac_ok}, {doac_fail}) | ((doac), {doac_fail}, {stop_doac}, {doac_danger})
| ((doac), {doac_danger}, {stop_doac}, {danger})

ADR  $\triangleq$  ((get_apixaban, get_diltiazem),  $\emptyset$ , {moderate}) | ((get_apixaban, diltiazem),  $\emptyset$ , {moderate}) | ((apixaban, get_diltiazem),  $\emptyset$ , {moderate})
| ((apixaban, diltiazem),  $\emptyset$ , {moderate}) | ((get_apixaban, get_verapamil),  $\emptyset$ , {moderate}) | ((get_apixaban, verapamil),  $\emptyset$ , {moderate})
| ((apixaban, get_verapamil),  $\emptyset$ , {moderate}) | ((apixaban, verapamil),  $\emptyset$ , {moderate}) | ((get_dabigatran, get_diltiazem),  $\emptyset$ , {moderate})
| ((get_dabigatran, diltiazem),  $\emptyset$ , {moderate}) | ((dabigatran, get_diltiazem),  $\emptyset$ , {moderate}) | ((dabigatran, diltiazem),  $\emptyset$ , {moderate})
| ((get_dabigatran, get_verapamil),  $\emptyset$ , {major}) | ((dabigatran, get_verapamil),  $\emptyset$ , {major}) | ((dabigatran, get_verapamil),  $\emptyset$ , {major})
| ((dabigatran, verapamil),  $\emptyset$ , {major}) | ((get_dabigatran, get_carvediol),  $\emptyset$ , {moderate}) | ((get_dabigatran, carvediol),  $\emptyset$ , {moderate})
| ((dabigatran, get_carvediol),  $\emptyset$ , {moderate}) | ((dabigatran, carvediol),  $\emptyset$ , {moderate}) | ((get_warfarin, get_benzepril),  $\emptyset$ , {minor})
| ((get_warfarin, benzepril),  $\emptyset$ , {minor}) | ((warfarin, get_benzepril),  $\emptyset$ , {minor}) | ((warfarin, benzepril),  $\emptyset$ , {minor})
| ((get_warfarin, get_indapamide),  $\emptyset$ , {minor}) | ((get_warfarin, indapamide),  $\emptyset$ , {minor}) | ((warfarin, get_indapamide),  $\emptyset$ , {minor})
| ((warfarin, indapamide),  $\emptyset$ , {minor}) | ((get_warfarin, get_chlorothiazide),  $\emptyset$ , {minor}) | ((get_warfarin, chlorothiazide),  $\emptyset$ , {minor})
| ((warfarin, get_chlorothiazide),  $\emptyset$ , {minor}) | ((warfarin, chlorothiazide),  $\emptyset$ , {minor}) | ((get_warfarin, get_propranolol),  $\emptyset$ , {minor})
| ((get_warfarin, propranolol),  $\emptyset$ , {minor}) | ((warfarin, get_propranolol),  $\emptyset$ , {minor}) | ((warfarin, propranolol),  $\emptyset$ , {minor})
| ((get_flecainide, get_diltiazem),  $\emptyset$ , {major}) | ((get_flecainide, diltiazem),  $\emptyset$ , {major}) | ((flecainide, get_diltiazem),  $\emptyset$ , {major})
| ((flecainide, diltiazem),  $\emptyset$ , {major}) | ((get_flecainide, get_verapamil),  $\emptyset$ , {major}) | ((get_flecainide, verapamil),  $\emptyset$ , {major})
| ((flecainide, get_verapamil),  $\emptyset$ , {major}) | ((flecainide, get_verapamil),  $\emptyset$ , {major}) | ((get_flecainide, get_bisoprolol),  $\emptyset$ , {moderate})
| ((get_flecainide, bisoprolol),  $\emptyset$ , {moderate}) | ((flecainide, get_bisoprolol),  $\emptyset$ , {moderate}) | ((flecainide, bisoprolol),  $\emptyset$ , {moderate})
| ((get_flecainide, get_atenolol),  $\emptyset$ , {moderate}) | ((get_flecainide, atenolol),  $\emptyset$ , {moderate}) | ((flecainide, get_atenolol),  $\emptyset$ , {moderate})
| ((flecainide, atenolol),  $\emptyset$ , {moderate}) | ((get_flecainide, get_propranolol),  $\emptyset$ , {moderate}) | ((get_flecainide, propranolol),  $\emptyset$ , {moderate})
| ((flecainide, get_propranolol),  $\emptyset$ , {moderate}) | ((flecainide, propranolol),  $\emptyset$ , {moderate}) | ((get_flecainide, get_carvediol),  $\emptyset$ , {moderate})
| ((get_flecainide, carvediol),  $\emptyset$ , {moderate}) | ((flecainide, get_carvediol),  $\emptyset$ , {moderate}) | ((flecainide, carvediol),  $\emptyset$ , {moderate})
| ((major),  $\emptyset$ , {major}) | ((moderate),  $\emptyset$ , {moderate}) | ((minor),  $\emptyset$ , {minor}) | ((alert_dup),  $\emptyset$ , {alert_dup}) | ((danger),  $\emptyset$ , {danger})

```

Fig. 20 Reactions for the comorbidity case study in 5.1.

1332 C.2 Auxiliary material for the comorbidity case study

The RS specification for the comorbidity case study presented in Bowles et al. (2024) is reported in Fig. 20 (set of reactions) and Fig. 21 (context processes definitions), where we assume the initial state is $D_0 = \emptyset$. The corresponding experimentation with GROOVE has been discussed in Sect. 5.1.

1333 C.3 Auxiliary material for the protein signaling networks case study

The BioResolve specification for the protein signaling networks case study presented in Ballis et al. (2024) is reported in Fig. 22. The corresponding experimentation with GROOVE has been discussed in Sect. 5.2.

1339

1340

1341

1342

1343

1344



```

eafib1  $\triangleq$   $(\emptyset, \{afib\}, \emptyset).eafib1 + ((afib), \emptyset, \emptyset).ehr$ 
ehr  $\triangleq$   $(\emptyset, \{heart\_rate\}, \emptyset).ehr + ([heart\_rate], \emptyset, \emptyset).ehr$ 
ebb  $\triangleq$   $\emptyset.ebb + e\_ebb + e\_cbb + e\_nsbb + e\_sbb$ 
e_cbb  $\triangleq$   $(\emptyset, \{verapamil\}, \{get\_diltiazem\}).empty + (\emptyset, \{diltiazem\}, \{get\_verapamil\}).empty$ 
e_nsbb  $\triangleq$   $(\emptyset, \{carvedilol\}, \{get\_propranolol\}).empty + (\emptyset, \{propranolol\}, \{get\_carvedilol\}).empty$ 
e_sbb  $\triangleq$   $(\emptyset, \{atenolol\}, \{get\_bisoprolol\}).empty + (\emptyset, \{bisoprolol\}, \{get\_atenolol\}).empty$ 
eafib2  $\triangleq$   $(\emptyset, \{afib\}, \emptyset).eafib2 + ((afib), \emptyset, \emptyset).ehf$ 
ehf  $\triangleq$   $(\emptyset, \{has\_fib\}, \emptyset).ehf + ([has\_fib], \emptyset, \emptyset).eflec$ 
eflec  $\triangleq$   $\emptyset.eflec + e\_flec$ 
e_flec  $\triangleq$   $(get\_flecainide).empty$ 
eafib3  $\triangleq$   $(\emptyset, \{afib\}, \emptyset).eafib3 + ((afib), \emptyset, \emptyset).econs$ 
econs  $\triangleq$   $(\emptyset, \{heart\_rate, has\_fib\}, \emptyset).econs + (\emptyset, \{consensus\_acei\}, \emptyset).econs + (\{consensus\_acei, heart\_rate\}, \emptyset, \emptyset).estroke + (\{consensus\_acei, has\_fib\}, \emptyset, \emptyset).estroke$ 
estroke  $\triangleq$   $(\emptyset, \{diseases, over75\}, \emptyset).ewarf + ([over75], \{doac\_fail, doac\_int\}, \emptyset).edoac + (\{diseases\}, \{doac\_fail, doac\_int\}, \emptyset).edoac + ([over75], \{doac\_fail\}, \emptyset, \emptyset).evkant$ 
+  $\triangleq$   $(\{over75, doac\_int\}, \emptyset, \emptyset).evkant + (\{diseases\}, \{doac\_fail\}, \emptyset, \emptyset).evkant$ 
ewarf  $\triangleq$   $\emptyset.ewarf + e\_warf$ 
e_warf  $\triangleq$   $(get\_warfarin).empty$ 
edoac  $\triangleq$   $\emptyset.edoac + e\_doac$ 
e_doac  $\triangleq$   $(\emptyset, \{dabigatran\}, \{get\_apixaban\}).e\_doacfail + (\emptyset, \{apixaban\}, \{get\_dabigatran\}).e\_doacfail$ 
e_doacfail  $\triangleq$   $(\{doac\_fail\}, \emptyset, \{stop\_doac\}).evkant + (\emptyset, \{doac\_fail\}, \emptyset).e\_doacfail$ 
evkant  $\triangleq$   $\emptyset.evkant + e\_vkant$ 
e_vkant  $\triangleq$   $(get\_vkant).empty$ 
ghyper  $\triangleq$   $(\emptyset, \{hyper\}, \emptyset).ghyper + ((hyper), \emptyset, \emptyset).g1$ 
g1  $\triangleq$   $(\{diabete\}, \emptyset, \emptyset).g2 + ([diabete, origin], \emptyset).g2 + (\emptyset, \{below55, diabete\}, \emptyset).g3 + (\{origin\}, \{diabete\}, \emptyset).g3$ 
g2  $\triangleq$   $\emptyset.g2 + (\emptyset, \{captoril\}, \{get\_benazepril\}).g4 + (\emptyset, \{benazepril\}, \{get\_captoril\}).g4 + (\emptyset, \{irbesartan\}, \{get\_olmesortan\}).g5$ 
g3  $\triangleq$   $\emptyset.g3 + (\emptyset, \{verapamil\}, \{get\_diltiazem\}).g6 + (\emptyset, \{diltiazem\}, \{get\_verapamil\}).g6$ 
g4  $\triangleq$   $\emptyset.g4 + (\emptyset, \{verapamil\}, \{get\_diltiazem\}).g7 + (\emptyset, \{diltiazem\}, \{get\_verapamil\}).g7 + (\emptyset, \{chlorothiazide\}, \{get\_indapamide\}).g8$ 
g5  $\triangleq$   $\emptyset.g5 + (\emptyset, \{verapamil\}, \{get\_chlorothiazide\}).g8$ 
g6  $\triangleq$   $\emptyset.g6 + (\emptyset, \{captoril\}, \{get\_benazepril\}).g9 + (\emptyset, \{benazepril\}, \{get\_captoril\}).g9 + (\emptyset, \{irbesartan\}, \{get\_olmesortan\}).g9$ 
g7  $\triangleq$   $\emptyset.g7 + (\emptyset, \{irbesartan\}, \{get\_olmesortan\}).g10 + (\emptyset, \{olmesortan\}, \{get\_irbesartan\}).g10$ 
g8  $\triangleq$   $\emptyset.g8 + (\emptyset, \{irbesartan\}, \{get\_olmesortan\}).ecbb + (\emptyset, \{olmesortan\}, \{get\_irbesartan\}).ecbb + (\emptyset, \{verapamil\}, \{get\_diltiazem\}).earb$ 
g9  $\triangleq$   $\emptyset.g9 + (\emptyset, \{captoril\}, \{get\_benazepril\}).ecbb + (\emptyset, \{benazepril\}, \{get\_captoril\}).ecbb + (\emptyset, \{chlorothiazide\}, \{get\_indapamide\}).earb$ 
g10  $\triangleq$   $\emptyset.g10 + (\emptyset, \{captoril\}, \{get\_benazepril\}).ecbb + (\emptyset, \{benazepril\}, \{get\_captoril\}).ecbb + (\emptyset, \{chlorothiazide\}, \{get\_indapamide\}).eacei$ 
g11  $\triangleq$   $\emptyset.g11 + (\emptyset, \{captoril\}, \{get\_benazepril\}).earb + (\emptyset, \{benazepril\}, \{get\_captoril\}).earb + (\emptyset, \{irbesartan\}, \{get\_olmesortan\}).eacei$ 
+  $\triangleq$   $(\emptyset, \{olmesortan\}, \{get\_irbesartan\}).eacei$ 
ecbb  $\triangleq$   $\emptyset.ecbb + e\_cbb$ 
eacei  $\triangleq$   $\emptyset.eacei + e\_acei$ 
e_acei  $\triangleq$   $(\emptyset, \{captoril\}, \{get\_benazepril\}).empty + (\emptyset, \{benazepril\}, \{get\_captoril\}).empty$ 
earb  $\triangleq$   $\emptyset.earb + e\_arb$ 
e_arb  $\triangleq$   $(\emptyset, \{irbesartan\}, \{get\_olmesortan\}).empty + (\emptyset, \{olmesortan\}, \{get\_irbesartan\}).empty$ 
etd  $\triangleq$   $\emptyset.etd + e\_td$ 
e_td  $\triangleq$   $(\emptyset, \{chlorothiazide\}, \{get\_indapamide\}).empty + (\emptyset, \{indapamide\}, \{get\_chlorothiazide\}).empty$ 
k_doac  $\triangleq$   $(\emptyset, \{doac\_test\}, \emptyset, \{doac\_ok\}).empty + ((doac\_test), \emptyset, \{doac\_fail\}).empty + (\emptyset, \{doac\_test\}, \emptyset).k\_doac$ 
empty  $\triangleq$   $\emptyset.empty$ 
kafib  $\triangleq$   $\{afib\}.empty + empty$ 
khf  $\triangleq$   $\{has\_fib\}.empty + empty$ 
khr  $\triangleq$   $\{heart\_rate\}.empty + empty$ 
kcons  $\triangleq$   $(\{consensus\_acei\}, empty + empty)$ 
kage  $\triangleq$   $(\{over75\}, empty + (\{below55\}, empty + empty))$ 
kdiabete  $\triangleq$   $(\{diabete\}, empty + empty)$ 
kdoacint  $\triangleq$   $(\{doac\_int\}, empty + empty)$ 
khyper  $\triangleq$   $(\{hyper\}, empty + empty)$ 
korigin  $\triangleq$   $(\{origin\}, empty + empty)$ 

```

Fig. 21 Context process definitions for the comorbidity case study in 5.1. The initial context is given by the parallel composition of therapies $eafib1 | eafib2 | eafib3 | ghyper$ and the parallel composition of features $kafib | khf | khr | kcons | kage | kdiabete | kdoacint | khyper | korigin | k_doac$.

1345 C.4 Auxiliary material for the T cell differentiation 1346 case study

1347 The BioResolve specification derived from the Boolean net-
1348 work model (available at Puniya (2024), see Fig. 24) of the
1349 T cell differentiation case study from Puniya et al. (2018),
1350 and exploited in Brodo et al. (2025) is reported in Fig. 24.
1351 The corresponding experimentation with GROOVE has been
1352 discussed in Sect. 5.3.

1353 C.5 Auxiliary material for the GROOVE experiments

1354 To replicate the GROOVE experiments reported in Sects. 4
1355 (for the toy running example) and 5, we have included the
1356 following supplementary resources with this submission:

- The start graphs derived from the BioResolve specifications in this appendix (C.1–C.4); 1358
- Instructions for calling the GROOVE generator so as to 1359 reproduce all the exploration runs, occurrence graphs and 1360 model checking results (using GROOVE version 7.4.3). 1361

- 1357 • The rule systems described in Sect. 4;

```

myentities([]).

myreactions([
    react([akt],[],[akt]),
    react([erbb3],[],[akt]),
    react([mtor],[],[akt]),
    react([pdk1],[],[akt]),
    react([erbb1],[e,p],[erbb1]),
    react([egf],[e,p],[erbb1]),
    react([plcg],[e,p],[erbb1]),
    react([erbb2],[e,t,p],[erbb2]),
    react([egf],[e,t,p],[erbb2]),
    react([erbb3],[e,t,p],[erbb2]),
    react([erbb3],[e,p],[erbb3]),
    react([hrg],[e,p],[erbb3]),
    react([erk12],[],[erk12]),
    react([egf],[],[erk12]),
    react([p],[],[erk12]),
    react([mek12],[],[mek12]),
    react([erbb1],[],[mek12]),
    react([erbb2],[],[mek12]),
    react([erbb3],[],[mek12]),
    react([mtor],[],[mtor]),
    react([p],[],[mtor]),
    react([akt],[],[mtor]),
    react([p70s6k],[],[p70s6k]),
    react([akt],[],[p70s6k]),
    react([mtor],[],[p70s6k]),
    react([erk12],[],[p70s6k]),
    react([pdk1],[],[pdk1]),
    react([erbb1],[],[pdk1]),
    react([erbb2],[],[pdk1]),
    react([erbb3],[],[pdk1]),
    react([mek12],[],[pdk1]),
    react([pkca],[],[pkca]),
    react([plcg],[],[pkca]),
    react([plcg],[],[plcg]),
    react([egf],[],[plcg]),
    react([erbb1],[],[plcg]),
    react([erbb2],[],[plcg]),
    react([erbb3],[],[plcg]) ]).

myenvironment("[
    k = {egf,hrg}.k,
    ket = {e,t}.ket,
    korep = ({e}.korep + {p}.korep),
    korept = ({e}.korept + {p}.korept + {t}.korept),
    kge = (?{erbb1},{}.{e}?.kge
        + ?{erbb2},{}.{e}?.kge
        + ?{},(erbb1,erbb2),{}?.kge) ]").

```

Fig. 22 BioResolve implementation of the protein signaling network case study from 5.2.

```

Jak1 = IFNgR and not SOCS1
IL21 = STAT3 and NFAT
IL18R = IL18 and IL12 and not STAT6
SOCS1 = STAT1 or Tbet
IL6 = RORgt
STAT5 = IL2R
IL17 = (RORgt and not STAT1)
or (STAT3 and IL17 and IL23R and not STAT1 and not STAT5)
STAT4 = IL12R and IL12 and not GATA3
IFNgR = (IFNg_e and NFAT) or (IFNg and NFAT)
STAT6 = IL4R and not IFNg and not SOCS1
GATA3 = (STAT6 and NFAT and not TGFB and not RORgt and not Foxp3
and not Tbet) or (GATA3 and not Tbet)
or (STAT5 and not TGFB and not RORgt and not Foxp3 and not Tbet)
IL4 = GATA3 and NFAT and not STAT1
NFKB = IRAK and not Foxp3
IL2 = NFAT and NFKB and not Tbet
IL23R = (IL23 and STAT3 and not Tbet) or STAT3
Tbet = (STAT4 and not RORgt and not Foxp3)
or (STAT1 and not RORgt and not Foxp3)
or (Tbet and not IL12 and not IFNg and not RORgt and not Foxp3)
TGFB = TGFB and NFAT
RORgt = TGFB and ((STAT3 and IL21R) or (STAT3 and IL6R)) and not Tbet
and not GATA3 and not Foxp3
IL6R = IL6 or IL6_e
IL21R = IL21
Foxp3 = (TGFB and not (IL6R and STAT3) and not IL21R and not GATA3)
or (STAT5 and not (IL6R and STAT3) and not IL21R and not GATA3)
IRAK = IL18R
IL12R = (IL12 and NFAT) or (STAT4 and not GATA3) or Tbet or (TCR and not GATA3)
IL2R = IL2 and NFAT
STAT3 = IL21R or IL23R or IL6R
IFNg = NFKB or (STAT4 and NFKB and NFAT and not STAT3 and not STAT6)
or (Tbet and not STAT3)
NFAT = TCR and not Foxp3
STAT1 = (IL27 and NFAT) or Jak1
IL4R = (IL4 and not SOCS1) or IL4_e

```

Fig. 23 Boolean updates of the T Cell differentiation model from Puniya et al. (2018), available at Puniya (2024).

Data Availability No datasets were generated or analysed during the current study. 1379
1380

Code Availability GROOVE is an open-source tool, available at <https://github.com/nl-utwente-groove/code>. For replication of the GROOVE experiments reported in this paper, we provide supplementary material described in Sect. C.5. 1381
1382
1383
1384

Declarations

Ethical approval Not applicable. 1386

Conflict of interest The authors declare no Conflict of interest. 1387

Open Access This article is licensed under a Creative Commons Attribution 4.0 International License, which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if changes were made. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by/4.0/>. 1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400

References

Ehrenfeucht A, Rozenberg G (2007) Reaction systems Fundam Informaticae 75(1–4):263–280 1401
1402
1403

```

myreactions([
  react([stat5],[gata3,i121r,i16r],[foxp3]),
  react([stat5],[gata3,i121r,stat3],[foxp3]),
  react([tgfb], [gata3,i121r,i16r],[foxp3]),
  react([tgfb], [gata3,i121r,stat3],[foxp3]),
  react([gata3],[tbet],[gata3]),
  react([nfat,stat6],[foxp3,rorgt,tbet,tgfb],[gata3]),
  react([nfat,nfkb,stat4],[stat3,stat6],[ifng]),
  react([nfkb],[],[ifng]),
  react([tbet],[stat3],[ifng]),
  react([ifng,nfat],[],[ifngr]),
  react([ifng,ifng],[nfat],[ifngr]),
  react([i112,nfat],[],[i112r]),
  react([stat4],[gata3],[i112r]),
  react([tbet],[],[i112r]),
  react([tcr],[gata3],[i112r]),
  react([i117,i123r,stat3],[stat1,stat5],[i117]),
  react([rorgt],[stat1],[i117]),
  react([i112,i118],[stat6],[i118r]),
  react([nfat,nfkb],[tbet],[i121]),
  react([nfat,stat3],[],[i121]),
  react([i121],[],[i121r]),
  react([i123,stat3],[tbet],[i123r]),
  react([stat3],[],[i123r]),
  react([i12,nfat],[],[i12r]),
  react([gata3,nfat],[stat1],[i14]),
  react([i14],[soc1],[i14r]),
  react([i14e],[],[i14r]),
  react([rorgt],[],[i16]),
  react([i16],[],[i16r]),
  react([i16e],[],[i16r]),
  react([i118r],[],[irak]),
  react([ifngr],[soc1],[jak1]),
  react([tcr],[foxp3],[nfat]),
  react([irak],[foxp3],[nfkb]),
  react([i121r,stat3,tgfb],[foxp3,gata3,tbet],[rorgt]),
  react([i16r,stat3,tgfb],[foxp3,gata3,tbet],[rorgt]),
  react([stat1],[],[soc1]),
  react([tbet],[],[soc1]),
  react([i127,nfat],[],[stat1]),
  react([jak1],[],[stat1]),
  react([i121r],[],[stat3]),
  react([i123r],[],[stat3]),
  react([i16r],[],[stat3]),
  react([i112,i112r],[gata3],[stat4]),
  react([i12r],[],[stat5]),
  react([i14r],[ifng,soc1],[stat6]),
  react([stat1],[foxp3,rorgt],[tbet]),
  react([stat4],[foxp3,rorgt],[tbet]),
  react([tbet],[foxp3,ifng,i112,rorgt],[tbet]),
  react([nfat,tgfb],[],[tgfb]) ]).

```

Fig. 24 BioResolve implementation of the T cell case study from 5.3.

ity treatment strategies. In: Gori, R., Milazzo, P., Tribastone, M. (eds.) Computational Methods in Systems Biology - 22nd International Conference, CMSB 2024, Pisa, Italy, September 16–18, Proceedings. Lecture Notes in Computer Science, vol. 14971, pp. 27–44. Springer, ??? . https://doi.org/10.1007/978-3-031-71671-3_3

Brodo L, Bruni R, Falaschi M, Gori R, Milazzo P (2025) Attractor and slicing analysis of a t cell differentiation model based on reaction systems. In: From Data to Models and Back - 11th International Symposium, DataMod 2023, Eindhoven, The Netherlands, November 6–7, 2023, Proceedings. Lecture Notes in Computer Science, vol. 14618, pp. 69–89. Springer, ??? . https://doi.org/10.1007/978-3-031-87217-4_4

Brodo L, Bruni R, Falaschi M (2021) A logical and graphical framework for reaction systems. *Theor Comput Sci* 875:1–27. <https://doi.org/10.1016/J.TCS.2021.03.024>

Ehrig H, Ehrig K, Prange U, Taentzer G (2006) Fundamentals of Algebraic Graph Transformation. Monographs in Theoretical Computer Science. An EATCS Series. Springer, ??? . <https://doi.org/10.1007/3-540-31188-2>

Heckel R, Taentzer G (2020) Graph Transformation for Software Engineers - With Applications to Model-Based Development and Domain-Specific Language Engineering. Springer, ??? . <https://doi.org/10.1007/978-3-030-43916-3>

Kreowski H, Rozenberg G (2019) Graph transformation through graph surfing in reaction systems. *J. Log. Algebraic Methods Program.* **109**. <https://doi.org/10.1016/J.JLAMP.2019.100481>

Ghamarian AH, Mol M, Rensink A, Zambon E, Zimakova M (2012) Modelling and analysis using GROOVE. *Int J Softw Tools Technol Transf* 14(1):15–40. <https://doi.org/10.1007/S10009-011-0186-X>

Ballis D, Brodo L, Falaschi M, Olarte C (2024) Process calculi and rewriting techniques for analyzing reaction systems. In: Gori, R., Milazzo, P., Tribastone, M. (eds.) Computational Methods in Systems Biology - 22nd International Conference, CMSB 2024, Pisa, Italy, September 16–18, 2024, Proceedings. Lecture Notes in Computer Science, vol. 14971, pp. 1–18. Springer, ??? . https://doi.org/10.1007/978-3-031-71671-3_1

Plotkin GD (2004) A structural approach to operational semantics. *J Log Algebraic Methods Program* 60–61:17–139

Milner R (1980) A Calculus of Communicating Systems. LNCS, vol. 92. Springer, ???

Rensink A (2024) The GROOVE tool set, https://github.com/nl-utwente-groove/code/releases/tag/release-7_0_1 version 7.0.1. Available at <https://github.com/nl-utwente-groove/code>

Brodo L, Bruni R, Falaschi M (2024) A framework for monitored dynamic slicing of reaction systems. *Nat Comput* 23(2):217–234. <https://doi.org/10.1007/S11047-024-09976-3>

Brodo L, Bruni R, Falaschi M, Gori R, Milazzo P, Montagna V, Pulieri P (2024) Causal analysis of positive reaction systems. *Int J Softw Tools Technol Transf* 26(4):509–526. <https://doi.org/10.1007/S10009-024-00757-Y>

Feder G, Eccles M, Grol R, Griffiths C, Grimshaw J (1999) Using clinical guidelines *Bmj* 318(7185):728–730. <https://doi.org/10.1136/bmj.318.7185.728>

Woolf SH, Grol R, Hutchinson A, Eccles M, Grimshaw J (1999) Potential benefits, limitations, and harms of clinical guidelines. *BMJ* 318(7182):527–530. <https://doi.org/10.1136/bmj.318.7182.527>

Hughes LD, McMurdo MET, Guthrie B (2013) Guidelines for people not for diseases: the challenges of applying UK clinical guidelines to people with multimorbidity. *Age and Ageing* 42, 62–69. <https://doi.org/10.1093/ageing/afs100>

Bowles JK, Caminati MB (2017) A flexible approach for finding optimal paths with minimal conflicts. In: Proceedings of ICFEM 2017. LNCS, vol. 10610, pp. 209–225. Springer, ??? . https://doi.org/10.1007/978-3-319-68690-5_13

- Azimi S, Iancu B, Petre I (2014) Reaction system models for the heat shock response. *Fund Inform* 131(3–4):299–312. <https://doi.org/10.3233/FI-2014-1016>
- Corollini L, Maj C, Marini F, Besozzi D, Mauri G (2012) An excursion in reaction systems: From computer science to biology. *Theor Comput Sci* 454:95–108. <https://doi.org/10.1016/j.tcs.2012.04.003>
- Azimi S (2017) Steady states of constrained reaction systems. *Theor Comput Sci* 701(C):20–26. <https://doi.org/10.1016/j.tcs.2017.03.047>
- Okubo F, Yokomori T (2016) The computational capability of chemical reaction automata. *Nat Comput* 15(2):215–224. <https://doi.org/10.1007/s11047-015-9504-7>
- Ehrenfeucht A, Main MG, Rozenberg G (2010) Combinatorics of life and death for reaction systems. *Int J Found Comput Sci* 21(3):345–356. <https://doi.org/10.1142/S0129054110007295>
- Ehrenfeucht A, Main MG, Rozenberg G (2011) Functions defined by reaction systems. *Int J Found Comput Sci* 22(1):167–178. <https://doi.org/10.1142/S0129054111007927>
- Bowles J, Brodo L, Bruni R, Falaschi M, Gori R, Milazzo P (2024) Enhancing reaction systems with guards for analysing comorbid-

- 1489 Clavel M, Durán F, Eker S, Lincoln P, Martí-Oliet N, Meseguer J, Talcott
 1490 CL (eds.) (2007) All About Maude - A High-Performance Logi-
 1491 cal Framework, How to Specify, Program and Verify Systems in
 1492 Rewriting Logic. Lecture Notes in Computer Science, vol. 4350.
 1493 Springer, ??. <https://doi.org/10.1007/978-3-540-71999-1>
- 1494 Heyde SV, Bender C, Henjes F (2014) Boolean erbB network reconstruc-
 1495 tions and perturbation simulations reveal individual drug response
 1496 in different breast cancer cell lines. *BMC Systems Biology* 8, 75.
<https://doi.org/10.1186/1752-0509-8-75>
- 1497 Rubio R, Martí-Oliet N, Pita I, Verdejo A (2021) Strategies, model
 1498 checking and branching-time properties in maude. *J Log Algebraic
 1499 Methods Program* 123:100700. <https://doi.org/10.1016/J.JLAMP.2021.100700>
- 1500 Brodo L, Bruni R, Falaschi M, Gori R, Milazzo P (2025) Slicing analy-
 1501 ses for negative dependencies in reaction systems modeling gene
 1502 regulatory networks. *Natural Computing*. Under review; will hope-
 1503 fully appear in same volume as this paper
- 1504 Saez-Rodriguez J, Simeoni L, Lindquist JA, Hemenway R, Bommhardt
 1505 U, Arndt B, Haus U-U, Weismantel R, Gilles ED, Klamt S (2007)
 1506 A logical model provides insights into T cell receptor signaling.
 1507 *PLoS Comput Biol* 3(8):163. <https://doi.org/10.1371/journal.pcbi.0030163>
- 1508 Thakar J, Albert R (2010) Boolean models of within-host immune inter-
 1509 actions. *Curr Opin Microbiol* 13(3):377–381. <https://doi.org/10.1016/j.mib.2010.04.003>
- 1510 Puniya BL, Todd RG, Mohammed A, Brown DM, Barberis M, Helikar T
 1511 (2018) A mechanistic computational model reveals that plasticity
 1512 of CD4+ T cell differentiation is a function of cytokine compo-
 1513 sition and dosage. *Frontiers in physiology* 9, 878. <https://doi.org/10.3389/fphys.2018.00878>
- 1514 Lafaille JJ (1998) The role of helper T cell subsets in autoimmune
 1515 diseases. *Cytokine & growth factor reviews* 9(2):139–151. [https://doi.org/10.1016/s1359-6101\(98\)00009-4](https://doi.org/10.1016/s1359-6101(98)00009-4)
- 1516 Hirahara K, Nakayama T (2016) CD4+ T-cell subsets in inflammatory
 1517 diseases: beyond the Th 1/Th 2 paradigm. *Int Immunopharmacol* 28(4):163–
 1518 171. <https://doi.org/10.1093/intimm/dxw006>
- 1519 Meng X, Yang J, Dong M, Zhang K, Tu E, Gao Q, Chen W, Zhang
 1520 C, Zhang Y (2016) Regulatory T cells in cardiovascular diseases.
 1521 *Nat Rev Cardiol* 13(3):167–179. <https://doi.org/10.1038/nrccardio.2015.169>
- 1522 Luckheeram RV, Zhou R, Verma AD, Xia B (2012) CD4+ T cells: differ-
 1523 entiation and functions. *Clinical and developmental immunology*
 1524 2012. <https://doi.org/10.1155/2012/925135>
- 1525 Helikar T, Kowal B, McClenathan S, Bruckner M, Rowley T, Madrahi-
 1526 mov A, Wicks B, Shrestha M, Limbu K, Rogers JA (2012) The cell
 1527 collective: toward an open and collaborative approach to systems
 1528 biology. *BMC Syst Biol* 6(1):1–14. <https://doi.org/10.1186/1752-0509-6-96>
- 1529 Azimi S, Gratie C, Ivanov S, Petre I (2015) Dependency graphs and mass
 1530 conservation in reaction systems. *Theor Comput Sci* 598:23–39.
<https://doi.org/10.1016/j.tcs.2015.02.014>
- 1531 Ivanov S, Rogojin V, Azimi S, Petre I (2018) WEBSIM: A web-
 1532 based reaction systems simulator. In: Díaz, C.G., Riscos-Núñez,
 1533 A., Paun, G., Rozenberg, G., Salomaa, A. (eds.) Enjoying Natural
 1534 Computing - Essays Dedicated to Mario de Jesús Pérez-Jiménez
 1535 on the Occasion of His 70th Birthday. Lecture Notes in Computer
 1536 Science, vol. 11270, pp. 170–181. Springer, ??. https://doi.org/10.1007/978-3-030-00265-7_14
- 1537 Nobile MS, Porreca AE, Spolaor S, Manzoni L, Cazzaniga P, Mauri
 1538 G, Besozzi D (2017) Efficient simulation of reaction systems on
 1539 graphics processing units. *Fundam Informaticae* 154(1–4):307–
 1540 321. <https://doi.org/10.3233/FI-2017-1568>
- 1541 Ferretti C, Leporati A, Manzoni L, Porreca AE (2020) The many roads
 1542 to the simulation of reaction systems. *Fundam Informaticae* 171(1–
 1543 4):175–188. <https://doi.org/10.3233/FI-2020-1878>
- 1544 Brodo L, Bruni R, Falaschi M, Gori R, Levi F, Milazzo P (2023) Quan-
 1545 titative extensions of reaction systems based on SOS semantics.
 1546 *Neural Comput Appl* 35(9):6335–6359. <https://doi.org/10.1007/S00521-022-07935-6>
- 1547 Meski A, Penczek W, Rozenberg G (2015) Model checking temporal
 1548 properties of reaction systems. *Inf. Sci.* 313, 22–42. <https://doi.org/10.1016/j.ins.2015.03.048>
- 1549 Meski A, Koutny M, Mikulski L, Penczek W (2024) Reaction mining
 1550 for reaction systems. *Nat Comput* 23(2):323–343. <https://doi.org/10.1007/S11047-024-09989-Y>
- 1551 Puniya BL (2024) CD4+ T cell Differentiation model webpage on the
 1552 CellCollective platform. Accessed: 18 March. <https://research.cellcollective.org/?dashboard=true#module/6678:1/cd4-t-cell-differentiation/1>
- 1553
- 1554
- 1555
- 1556
- 1557
- 1558
- 1559
- 1560
- 1561
- 1562
- 1563
- 1564
- 1565
- 1566
- 1567
- 1568
- 1569

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.